

# Page Replacement Policies

## PAGE REPLACEMENT — NAIVE

### step 1: save/clear victim page:

- drop page if fetched from disk and clean
- *dirty*: write back modifications if from disk and dirty (unless `MAP_COPY`)
- *non-dirty*: write page file/swap partition otherwise (e.g., stack, heap memory)

### step 2: unmap page from old AS: invalidate PTE, flush cache

### step 3: prepare new page: null page or load new contents

### step 4: map page frame into new AS: invalidate PTE, flush cache

## PAGE REPLACEMENT — BUFFERING

**problem:** naive page replacement encompasses two I/O transfers

- both operations block page fault from completing

**goal:** reduce I/O from critical page fault path to speed up page faults

**idea:** keep pool of free page frames (*pre-cleaning*):

- *on page fault*: use page frame from free pool
- *cleaning*: daemon cleans, reclaims and scrubs pages for free pool in background
- smooths out I/O, speeds up paging significantly

**remaining problem:** which pages to select as victims?

- *goal*: identify page that has left working set of its processes, add to free pool
- *success metric*: low overall page fault rate

## PAGE REPLACEMENT — FIFO

**idea:** evict oldest fetched page in system

**Belady's Anomaly:** using FIFO, for every number  $n$  of page frames you can construct a reference string that performs worse with  $n + 1$  frames

- with FIFO it is possible to get more page faults with more page frames!

## PAGE REPLACEMENT — ORACLE

= optimal replacement strategy: replace page whose next reference is furthest in future

**problem:** future unpredictable

**however:** good metric to check how well other algorithms perform

## PAGE REPLACEMENT — LRU

**goal:** approximate oracle page replacement

**idea:** past often predicts future well

**assumption:** page used furthest in past is used furthest in future

**cycle counter implementation:**

- have MMU write CPU's time stamp counter to PTE on every access
- *page fault*: scan all PTEs to find oldest counter value
- *advantage*: cheap at access if done in HW
- *disadvantage*: memory traffic for scanning

**stack implementation:**

- keep doubly linked list of all page frames
- move each referenced page to tail of list
- *advantage*: can find replacement victim in  $O(1)$
- *disadvantage*: need to change 6 pointers at every access

~> **no silver bullet:**

- *observation*: predicting future based on past is not precise
- *conclusion*: relax requirements — maybe perfect LRU isn't needed? ⇒ approximate LRU

## LRU APPROXIMATION — CLOCK PAGE REPLACEMENT

aka *second chance page replacement*

**precondition:** MMU sets reference bit in PTE

- supported natively by most hardware
- can easily emulate in systems with software managed TLB (e.g., MIPS)

**store:** keep all pages in circular FIFO list

**searching** for victim: scan pages in FIFO's order

- if reference bit = 0 → use page as victim and advance
- if reference bit = 1 → set to 0, continue scanning

**problem:** large memory → most pages referenced before scanned

- *solution*: use 2 arms, leading arm clears reference bit, trailing arm selects victim

## REPLACEMENT STRATEGIES — OTHER

**random eviction:** pick random victim

- dirt simple
- not overly horrible in reality

**larger counter:** use  $n$ -bit reference counter instead of reference bit

- *least frequently used*: rarely used page not in a working set → replace page with smallest count
- *most frequently used*: page with smallest count probably just brought in → replace page with largest count
- neither LFU nor MDU are common (no such hardware, not that great)

### Summary

victim page frame needs to be selected by OS when handling page faults

- evicting page frame after page fault happens = not a good idea
- page buffering keeps eviction out of critical path

different victim selection policies exist

- FIFO → Belady's Anomaly
- Oracle → cannot predict the future
- Random → unpredictable, never great but rarely very bad
- LRU → hard to implement efficiently