

# I/O Systems

## DEVICE MANAGEMENT — OBJECTIVES

**abstraction** from details of physical devices

**uniform naming** that does not depend on hardware details

**serialization** of I/O operations by concurrent applications

**protection** of standard-devices against unauthorized accesses

**buffering** if data from/to device cannot be stored in final destination

**error handling** of sporadic device errors

**virtualizing** physical devices via memory + time multiplexing

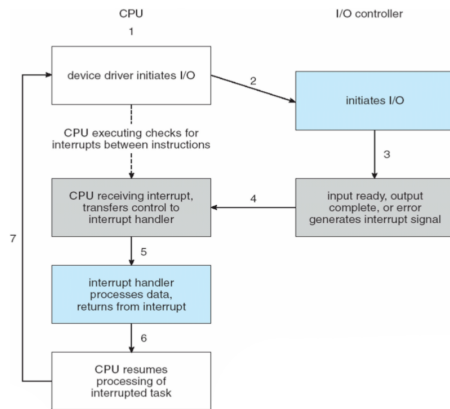
## DEVICE MANAGEMENT — TECHNIQUES

### programmed I/O:

- thread is busy-waiting for I/O operation to complete → CPU cannot be used elsewhere
- kernel is *polling* state of I/O device (command-ready, busy, error)

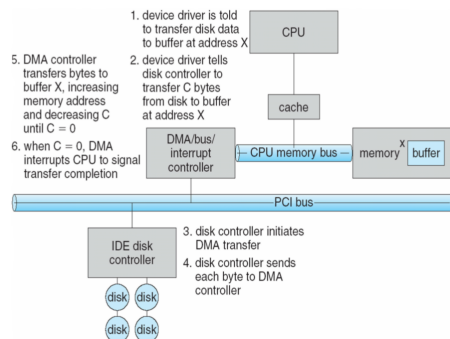
### interrupt-driven I/O:

- I/O command is issued
- processor continues executing instructions
- I/O device sends interrupt when command is done



### direct memory access (DMA):

- DMA module controls exchange of data between main memory and I/O device
- processor interrupted after entire block has been transferred
- bypasses CPU to transfer data directly between I/O device and memory



## KERNEL I/O SUBSYSTEM

**scheduling:** order I/O requests in per-device queues

**buffering:** store data in memory while transferring between devices

**error handling:** recover from read/availability/write errors

**protection:** protect from accidental/purposeful disruptions

**spooling:** hold output to device if device is slow (e.g., printer)

**reservation:** provide exclusive access for process

## DEVICE DRIVERS

### jobs:

- *translate* user request through device-independent standard interface
- *initialize* hardware at boot time
- *shut down* hardware

## DEVICE BUFFERING

### reasons:

- without buffering threads must wait for I/O to complete before proceeding
- pages must remain in main memory during physical I/O

### version 1 — block-oriented:

- information is stored in fixed-size blocks
- transfers are made a block at a time
- used for disks/tapes

### version 2 — stream-oriented:

- transfer information as byte stream
- used for keyboard, terminals, ... (most things that is not secondary storage)

## BUFFERING — USER LEVEL

**principle:** task specifies memory buffer where incoming data is placed

### issues:

- what happens if buffer is currently paged out to disk? → data loss
- additional problems with writing? → when is buffer available for re-use?

## BUFFERING — SINGLE

**principle:** user process can process one data block while next block is read in

**swapping:** can occur since input is taking place in system memory, not user memory

**stream-oriented:** buffer = input line, carriage return signals end of line

### block-oriented:

- input transfers made to *system buffer*
- buffer moved to *user space* when needed
- another block read into system buffer

## BUFFERING — DOUBLE

**principle:** use 2 system buffers instead of 1 (per user process)

user process can write/read from one buffer while OS empties/fills other buffer

## BUFFERING — CIRCULAR

**problem:** double buffer insufficient for high-burst traffic situations:

- many writes between long periods of computations
- long computation periods while receiving data
- might want to read ahead more than just single block from disk