

# Intro to C

## HELLO WORLD

```
#include <stdio.h>

int main(void) {
    printf("Hello World!\n");
    return 0;
}
```

`#include`: preprocessor inserts `stdio.h` contents  
`stdio.h`: contains `printf` declaration  
`main`: program starts here  
`void`: keyword for argument absence  
`{ }`: basic block/scope delimiters  
`printf`: prints to the terminal  
`\n`: newline character  
`return`: leave function, return value

## COMPILING

```
$ gcc hello.c -o hello
$ ./hello
Hello World!
```

## BASIC DATA TYPES

`char c = 5; char c = 'a';`  
 one byte, usually for characters (1970: ASCII is fine)  
`int i = 5; int i = 0xf; int i = 'a';`  
 usually 4 bytes, holds integers  
`float f = 5; float f = 5.5;`  
 4 bytes floating point number  
`double d = 5.19562`  
 8 bytes double precision floating point number

## BASIC DATA TYPES — LOGIC

```
int i = 5 / 2; //i = 2
    integer logic, no rounding
float f = 5.0f / 2; //f = 2.5f
    decimal logic for float and double
char a = 'a' / 2 //a = 97 / 2 = 48
    char interpreted as character by console
```

## BASIC DATA TYPES — SIGNED/UNSIGNED

```
signed int i = -5 //i = -5 (two's complement)
unsigned int i = -5 //i = 4294967291
```

## BASIC DATA TYPES — SHORT/LONG

```
short int i = 1024 //-32768...32767
long int i = 1024 //-2147483648...2147483647
```

## BASIC DATA TYPES — MORE SIZE STUFF

`sizeof int; sizeof long int; //4; 4; (x86 32-Bit)`  
 use data types from `inttypes.h` to be sure about sizes:

```
#include <inttypes.h>
int8_t i; uint32_t j;
```

## BASIC DATA TYPES — CONST/VOLATILE

```
const int c = 5;
    i is constant, changing it will raise compiler error
volatile int i = 5;
    i is volatile, may be modified elsewhere (by different program in shared memory, important for
    CPU caches, register, assumptions thereof)
```

## VARIABLES — LOCAL VS. GLOBAL

```
int m; // global variable

int myroutine(int j) {
    int i = 5 // local variable
    i = i+j;
    return i;
}
```

global variables (`int m`):

- lifetime: while program runs
- placed on pre-defined place in memory

basic block/function-local variables (`int i`):

- lifetime: during invocation of routine
- placed on stack or in registers

## VARIABLES — LOCAL VS. STATIC

```
int myroutine(int j) {
    static int i = 5;
    i = i+j;
    return i;
}

k = myroutine(1); // k = 6
k = myroutine(1); // k = 7
```

static function-local variables:

- saved like global variables
- variable persistent across invocations
- lifetime: like global variables

## PRINTING

```
int i = 5; float f = 2.5;
printf("The numbers are i=%d, f=%f", i, f);
```

comprised of format string and arguments

may contain format identifiers (`%d`)

see also [man printf](#)

special characters: encoded via leading backslash:

```
\n newline
\t tab
\' single quote
\" double quote
\0 null, end of string
```

## COMPOUND DATA TYPES

**structure**: collection of named variables (different types)

**union**: single variable that can have multiple types

members accessed via `.` operator

```
struct coordinate {
    int x;
    int y;
}
```

```

union longorfloat {
    long l;
    float f;
}

struct coordinate c;
c.x = 5;
c.y = 6;

union longorfloat lf;
lf.l = 5;
lf.f = 6.192;

```

## FUNCTIONS

encapsulate functionality (*reuse*)

code structuring (*reduce complexity*)

must be **declared** and **defined**

Declaration: states signature

Definition: states implementation (implicitly declares function)

```

int sum(int a, int b); // declaration

int sum(int a, int b) { // definition
    return a+b;
}

```

## HEADER FILES

header file for frequently used declarations

use **extern** to declare global variables defined elsewhere

use **static** to limit scope to current file (e.g. `static float pi` in `sum.c`: no `pi` in `main.c`)

```

// mymath.h
int sum(int a, int b);
extern float pi;

// sum.c
#include "mymath.h"

float pi = 3.1415927;
int sum(int a, int b) {
    return a+b;
}

// main.c
#include <stdio.h>
#include "mymath.h"

void main() {
    printf("%d\n", sum(1,2));
    printf("%f\n", pi);
}

```

## DATA SEGMENTS AND VARIABLES

Stack: local variables

Heap: variables created at runtime via `malloc()`/`free()`

Data Segment: static/global variables

Code: functions

## FUNCTION OVERLOADING

no function overloading in C!

use arrays over pointers

## POINTERS

```

int a = 5;
int *p = &a // points to int, initialized to point to a
int *q = 32 // points to int at address 32
int b = a+1;
int c = *p; // dereference(p) = dereference(&a) = 5
int d = (*p)+2 // = 7
int *r = p+1; // pointing to next element p is pointing to
int e = *(p+2) // dereference (p+2) = d = 7

```

## POINTERS — LINKED LIST

linked-list implementation via next-pointer

```

struct ll {
    int item;
    struct ll *next;
}

struct ll first;
first.item = 123;

struct ll second;
second.item = 456;
first.next = &second;

```

## ARRAYS

= fixed number of variables *continuously laid out in memory*

```

int A[5]; // declare array (reserve memory space)
A[4] = 25; A[0] = 24; // assign 25 to last, 24 to first elem
char c[] = {'a',5,6,7,'B'} // init array, length implicit
c[64] = 'Z' // NO bounds checking at compile/run (may raise protection fault)

// declare pointer to array; address elements via pointer:
char *p = c;
*(p+1) = 'Z'; p[3] = 'B'; char b = *p; // = 'a'

```

## STRINGS

= array of **chars** terminated by **NULL**:

```

char A[] = { 'T', 'e', 's', 't', '\0' };
char A[] = "Test";

```

declaration via pointer:

```
const char *p = "Test";
```

common string functions (`string.h`):

length: `size_t strlen(const char *s, size_t maxlen)`

compare:

```
int strcmp(const char *s1, const char *s2, size_t n);
```

copy: `int strncpy(char *dest, const char *src, size_t n);`

tokenize: `char *strtok(char *str, const char *delim);`

(e.g. split line into words)

## ARITHMETIC/BITWISE OPERATORS

arithmetic operators:

```
a+b, a++, ++a, a+=b, a-b, a--, --a, a-=b, a*b, a*=b, a/b, a/=b, a%b, a%=b
```

logical operators:

```
a&b, a|b, a>>b, a<<b, a^b, ~a
```

difference pre-/post-increment:

```
int a = 5;
if(a++ == 5) printf("Yes"); // Yes
a = 5;
if(++a == 5) printf("Yes"); // nothing
```

operators in order of precedence:

```
() , [] , -> , .
!, ++, --, +y, -y, *z, &=, (type), sizeof
*, /, %
+, -
<<, >>
<, <=, >, >=
==, !=
&
~
|
&&
||
?, :
=, +=, -=, *=, /=, %=, &=, ~=, <=, >=|
,
```

## STRUCTURES

brackets only needed for multiple statements

if/else, for, while, do-while, switch

may use break/continue

switch: need break statement, otherwise will fall through

```
if(a==b) printf("Equal") else printf("Different");
for(i=10; i>=10; i--) printf("%d", i+1);
int i=10; while(i--) printf("foo");
int i=0; do printf("bar"); while(i++ != 0);
```

```
char a = read();
switch(a) {
    case '1':
        handle_1();
        break;
    default:
        handle_other();
        break;
}
```

## TYPE CASTING

explicit casting: precision loss possible

```
int i = 5; float f = (float)i;
```

implicit casting: if no precision is lost

```
char c = 5; int i = c;
```

pointer casting: changes address calculation

```
int i = 5; char *p = (char *)&i; *(p+1) = 5;
```

type hierarchy: „wider“/„shorter“ types

```
unsigned int wider than signed int
```

operators cast parameters to widest type

Attention: assignment cast after operator cast

## C PREPROCESSOR

modifies source code before compilation

based on preprocessor *directives* (usually starting with #)

```
#include <stdio.h>, #include "mystdio.h":
```

copies contents of file to current file

only works with strings in source file

completely ignores C semantics

## PREPROCESSOR — SEARCH PATHS

#include <file>: system include, searches in:

```
/usr/local/include
libdir/gcc/[target]/[version]/include
/usr/[target]/include
/usr/include
(target: arch-specific (e.g. i686-linux-gnu),
version: gcc version (e.g. 4.2.4))
```

#include "file": local include, searches in:

```
directory containing current file
then paths specified by -i <dir>
then in system include paths
```

## PREPROCESSOR — DEFINITIONS

defines introduce replacement strings (can have arguments, based on string replacement)  
can help code structuring, often leading to source code cluttering

```
#define PI 3.14159265
#define TRUE (1)
#define max(a,b) ((a > b) ? (a) (b))
#define panic(str) do { printf(str); for (;;) } while(0);

#ifdef __unix__
# include <unistd.h>
#elif defined _WIN32
# include <windows.h>
#endif
```

## PREPROCESSOR — PREDEFINED MACROS

system-specific:

```
__unix__, _WIN32, __STDC_VERSION__
```

useful:

```
__LINE__, __FILE__, __DATE__
```

## LIBRARIES

= collection of functions contained in object files, glued together in dynamic/static library

ex.: Math header contains declarations, but not all definitions

→ need to link math library: `gcc math.c -o math -lm`

```
#include <math.h>
#include <stdio.h>
```

```
int main() {
    float f = 0.555f;
    printf("%f", sqrt(f*4));
    return 0;
}
```

# Introduction to Operating Systems

## WHAT'S AN OS?

**abstraction:** provides abstraction for applications

manages and hides hardware details

uses low-level interfaces (not available to applications)

multiplexes hardware to multiple programs (*virtualisation*)

makes hardware use efficient for applications

**protection:**

from processes using up all resources (*accounting, allocation*)

from processes writing into other processes memory

**resource managing:**

manages + multiplexes hardware resources

decides between conflicting requests for resource use

strives for efficient + fair resource use

**control:**

controls program execution

prevents errors and improper computer use


→ **no universally accepted definition**

**HARDWARE OVERVIEW**

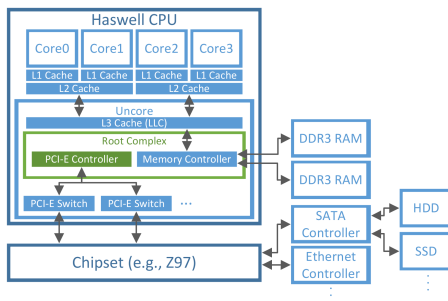
CPU(s)/devices/memory (conceptually) connected to common bus

CPU(s)/devices competing for memory cycles/bus

all entities run concurrently



today: multiple busses

**CENTRAL PROCESSING UNIT (CPU) — OPERATION**

fetches instructions from memory, executes them (instruction format/-set depends on CPU)

CPU internal registers store (meta-)data during execution (general purpose registers, floating point

registers, instruction pointer (IP), stack pointer (SP), program status word (PSW),...)

**execution modes:**

**user mode** (x86: *Ring 3/CPL 3*):


only non-privileged instructions may be executed

cannot manage hardware → **protection**

**kernel mode** (x86: *Ring 0/CPL 0*):

all instructions allowed

can manage hw with **privileged instructions**

**RANDOM ACCESS MEMORY (RAM)**

keeps currently executed instructions + data

today: CPUs have built-in *memory controller*

root complex connected directly via

„Wire“ to caches

pins to RAM

pins to PCI-E switches

**CACHING**

RAM delivers instructions/data slower than CPU can execute

memory references typically follow *locality principle*:

**spatial locality**: future refs often near previous accesses

(e.g. next byte in array)

**temporal locality**: future refs often at previously accessed ref

(e.g. loop counter)

*caching* helps mitigating this memory wall:

copy used information temporarily from slower to faster storage

check faster storage first before going down **memory hierarchy**

if not, data is copied to cache and used from there

**Access latency:**

register: ~1 CPU cycle

L1 cache (per core): ~4 CPU cycles

L2 cache (per core pair): ~12 CPU cycles

L3 cache/LLC (per uncore): ~28 CPU cycles (~25 GiB/s)

DDR3-12800U RAM: ~28 CPU cycles + ~50ns (~12 GiB/s)

**CACHING — CACHE ORGANISATION**

caches managed in hardware

divided into *cache lines* (usually 64 bytes each, unit at which data is exchanged between hierarchy levels)

often separation of data/instructions in faster caches (e.g. L1, see *harward architecture*)

**cache hit**: accessed data already in cache (e.g. L2 cache hit)

**cache miss**: accessed data has to be fetched from lower level

cache miss types:

*compulsory miss*: first ref miss, data never been accessed

*capacity miss*: cache not large enough for process working set

*conflict miss*: cache has still space, but collisions due to placement strategy

**INTERPLAY OF CPU AND DEVICES**

I/O devices and CPU execute concurrently


Each device controller

- is in charge of particular device

- has local buffer

**Workflow:**

1. CPU issues commands, moves data to devices
2. Device controller informs APIC (*Advanced Programmable Interrupt Controller*) that operation is finished
3. APIC signals CPU
4. CPU receives device/interrupt number from APIC, executes handler



## DEVICE CONTROL

Devices controlled through their **device controller**, accepts commands from OS via **device driver**  
 devices controlled through device registers and device memory:

*control* device by writing device registers

*read* status of device by reading device registers

*pass data* to device by reading/writing device memory

2 ways to access device registers/memory:


1. **port-mapped IO (PMIO):**  
 use special CPU instructions to access port-mapped registers/memory  
 e.g. x86 has different **in/out**-commands that transfer 1,2 or 4 bytes between CPU and device
2. **memory-mapped IO (MMIO):**  
 use same address space for RAM and device memory  
 some addresses map to RAM, others to different devices  
 access device's memory region to access device registers/memory

some devices use hybrid approaches using both

## DEVICE CONTROL — NVIDIA GENERAL PURPOSE GPU

memory-mapped ring-buffer and **put/get**-device

mapping can be exposed to application → application can submit commands in user-mode



### Summary

The OS is an abstraction layer between applications and hardware (multiplexes hardware, hides hardware details, provides protection between processes/users)

The CPU provides a separation of User and Kernel mode (which are required for an OS to provide protection between applications)

CPU can execute commands faster than memory can deliver instructions/data – memory hierarchy mitigates this memory wall, needs to be carefully managed by OS to minimize slowdowns

device drivers control hardware devices through PMIO/MMIO

Devices can signal the CPU (and through the CPU notify the OS) through interrupts

## OS Concepts

### OS INVOKATION

OS Kernel does **not** always run in background!

Occasions invoking kernel, switching to kernel mode:

1. **System calls:** User-Mode processes require higher privileges
2. **Interrupts:** CPU-external device sends signal
3. **Exceptions:** CPU signals unexpected condition

## SYSTEM CALLS — MOTIVATION

**Problem:** protect processes from one another

**Idea:** Restrict processes by running them in user-mode

→ **Problem:** now processes cannot manage hardware,...

who can switch between processes?

who decides if process may open certain file?

→ **Idea:** OS provides **services** to apps

app calls system if service is needed (**syscall**)

OS checks if app is allowed to perform action

if app may perform action and hasn't exceeded quota,

OS performs action in behalf of app in kernel mode

## SYSTEM CALLS — EXAMPLES

`fd = open(file, how, ...)` – open file for read/write/both

documented e.g. in **man 2 write**

overview in **man 2 syscalls**

## SYSTEM CALLS VS. APIS

**syscalls:** interface between apps and OS services, limited number of well-defined entry points to kernel

**APIs:** often used by programmers to make syscalls

e.g. **printf** library call uses **write** syscall

common APIs: Win32, POSIX, C API

## SYSTEM CALLS — IMPLEMENTATION

**trap instruction:** single syscall interface (entry point) to kernel

switches CPU to kernel mode, enters kernel in same, predefined

way for all syscalls

system call dispatcher then acts as syscall multiplexer

syscalls identified by number passed to trap instruction

**syscall table** maps syscall numbers to kernel functions

dispatcher decides where to jump based on number and table

programs (e.g. **stdlib**) have syscall number compiled in!

→ never reuse old numbers in future kernel versions

## INTERRUPTS

devices use interrupts to signal predefined conditions to OS

reminder: device has „interrupt line“ to CPU

e.g. device controller informs CPU that operation is finished

**programmable interrupt controller** manages interrupts

interrupts can be **masked**

masked interrupts: queued, delivered when interrupt unmasked

queue has finite length → interrupts can get lost

noteable interrupt examples:

1. **timer-interrupt:** periodically interrupts processes, switches to kernel → can then switch to different processes for fairness
2. **network interface card** interrupts CPU when packet was received → can deliver packet to process and free NIC buffer

when interrupted, CPU

1. looks up **interrupt vector** (= table pinned in memory, contains addresses of all service routines)
2. transfers control to respective **interrupt service routine** in OS that handles interrupt

interrupt service routine must first save interrupted process's state (instruction pointer, stack pointer, status word)

## EXCEPTIONS

sometimes unusual condition makes it impossible for CPU to continue processing

→ **Exception** generated within CPU:

1. CPU interrupts program, gives kernel control
2. kernel determines reason for exception
3. if kernel can resolve problem → does so, continues **faulting instruction**
4. kills process if not

Difference to Interrupts: interrupts can happen in any context, exceptions always occur asynchronous and in process context

## OS CONCEPTS — PHYSICAL MEMORY

up to early 60s:

- programs loaded and run directly in *physical memory*
- program too large → partitioned manually into *overlays*
- OS then swaps overlays between disk and memory
- different jobs could observe/modify each other

## OS CONCEPTS — ADDRESS SPACES

bad programs/people need to be isolated

**Idea:** give every job the illusion of having all memory to itself  
every job has own *address space*, can't name addresses of others  
jobs always and only use virtual addresses

## VIRTUAL MEMORY — INDIRECT ADDRESSING

Today: every CPU has built-in **memory management unit (MMU)**

MMU translates virtual addresses to physical addresses at every store/load operation

→ address translation protects one program from another

Definitions:

**Virtual address:** address in process' address space

**Physical address:** address of real memory

## VIRTUAL MEMORY — MEMORY PROTECTION

MMU allows kernel-only virtual addresses

- kernel typically part of all address spaces
- ensures that apps can't touch kernel memory

MMU can enforce *read-only* virtual addresses

- allows safe sharing of memory between apps

MMU can enforce execute disable

- makes code injection attacks harder

## VIRTUAL MEMORY — PAGE FAULTS

not all addresses need to be mapped at all times

- MMU issues *page fault* exception when accessed virtual address isn't mapped
- OS handles page faults by loading faulting addresses and then continuing the program
- → memory can be **over-committed**: more memory than physically available can be allocated to application

page faults also issued by MMU on illegal memory accesses

## OS CONCEPTS — PROCESSES

= program in execution („instance“ of program)

each process is associated with a **process control block (PCB)**

contains information about allocated resources

each process is associated with a virtual **address space (AS)**

- all (virtual) memory locations a program can name
- starts at 0 and runs up to a maximum
- address 123 in AS1 generally ≠ address 123 in AS2
- indirect addressing → different ASes to different programs
- → protection between processes

## OS CONCEPTS — ADDRESS SPACE LAYOUT


address spaces typically laid-out in different **sections**

- memory addresses between sections **illegal**
- illegal addresses → page fault
- more specifically calls **segmentation fault**
- OS usually kills process causing segmentation fault

**Stack:** function history, local variables

**Data:** Constants, static/global variables, strings

**Text:** Program code



## OS CONCEPTS — THREADS

each process:  $\geq 1$  threads (representing execution states)

IP stores currently executed instruction (address in **text** section)


SP register stores address of stack top

( $> 1$  threads → multiple stacks!)

PSW contains flags about execution history

(e.g. last calculation was 0 → used in following jump instruction)

more general purpose registers, floating point registers,...



## OS CONCEPTS — POLICIES VS. MECHANISMS

separation useful when designing OS

**Mechanism:** implementation of what is done

(e.g. commands to put a HDD into standby mode)

**Policy:** rules which decide when what is done and how much

(e.g. how often, how many resources are used,...)

→ *mechanisms can be reused even when policy changes*

## OS CONCEPTS — SCHEDULING


multiple processes/threads available → OS needs to switch between them (for multitasking)

*scheduler* decides which job to run next (policy)

*dispatcher* performs task-switching (mechanism)

schedulers try to

- provide fairness
- while meeting goals
- and adhering to priorities



## OS CONCEPTS — FILES

OS hides peculiarities of disks,...

programmer uses device-independent *files/directories* for persistent storage

**Files:** associate *file name* and *offset* with bytes

**Directories:** associate *directory names* with directory names or file names

**File System:** ordered block collection

- main task: translate (dir name + file name + offset) to block

- programmer uses file system operations to operate on files

(*open, read, seek*)

processes can communicate directly through special *named pipe* file (used with same operations as any other file)

## OS CONCEPTS — DIRECTORY TREE

directories form *directory tree/file hierarchy*  
 → structure data  
*root directory*: topmost directory in tree  
 files specified by providing *path name* to file

## OS CONCEPTS — MOUNTING

Unix: common to orchestrate multiple file systems in single file hierarchy  
 file systems can be *mounted* on directory  
 Win: manage multiple directory hierarchies with drive letters  
 (e.g. C: \Users)

## OS CONCEPTS — STORAGE MANAGEMENT

OS provides uniform view of information storage to file systems

- *drivers* hide specific hardware devices  
 → hides device peculiarities
- general interface abstracts physical properties to logical units  
 → block

OS increases I/O performance:

- **Buffering**: Store data temporarily while transferred
- **Caching**: Store data parts in faster storage
- **Spooling**: Overlap one job's output with other job's input

# Processes

## THE PROCESS ABSTRACTION

computers do „several things at the same time“ (just looks this way though quick process switching  
 (**Multiprogramming**))

→ **process** abstraction models this concurrency:

- container contains information about program execution
- conceptually, every process has own „virtual CPU“
- execution context is changed on process switch
- dispatcher switches context when switching processes
- **context switch**: dispatcher saves current registers/memory mappings, restores those of next process

## PROCESS-COOKING ANALOGON

Program/Process like Recipe/Cooking

**Recipe**: lists ingredients, gives algorithm what to do when  
 → program describes memory layout/CPU instructions

**Cooking**: activity of using the recipe  
 → process is activity of executing a program

multiple similar recipes for same dish  
 → multiple programs may solve same problem

recipe can be cooked in different kitchens at the same time  
 → program can be run on different CPUs at the same time  
 (as different processes)

multiple people can cook one recipe  
 → one process can have several worker threads

## CONCURRENCY VS. PARALLELISM

OS uses currency + parallelism to implement multiprogramming

1. **Concurrency**: multiple processes, one CPU  
 → not at the same time
2. **Parallelism**: multiple processes, multiple CPU  
 → at the same time

## VIRTUAL MEMORY ABSTRACTION — ADDRESS SPACES

every process has own *virtual addressess* (vaddr)  
 MMU relocates each load/store to *physical memory* (pmem)  
 processes never see physical memory, can't access it directly  
 + MMU can enforce protection (mappings in kernel mode)  
 + programs can see more memory than available  
   80:20 rule: 80% of process memory idle, 20% active  
   can keep working set in RAM, rest on disk  
 - need special MMU hardware

## ADDRESS SPACE (PROCESS VIEW)

code/data/state need to be organized within process  
 → **address space layout**  
 Data types:

1. *fixed size* data items
2. data naturally *free'd* in *reverse allocation order*
3. data *allocated/free'd „randomly“*

compiler/architecture determine how large int is and what instructions are used in text section (**code**)

**Loader** determines based on exe file how executed program is placed in memory

## SEGMENTS — FIXED-SIZE DATA + CODE

some data in programs never changes or will be written but never grows/shrinks  
 → memory can be statically allocated on process creation

**BSS segment** (*block started by symbol*):

- statically allocated variables/non-initialized variables
- executable file typically contains starting address + size of BSS
- entire segment initially 0

**Data segment**:

- fixed-size, initlized data elements (e.g. global variables)

**read-only data segment**:

- constant numbers, strings

All three sometimes summarized as one segment

compiler and OS decide ultimately where to place which data/how many segments exist

## SEGMENTS — STACK

some data naturally free'd in reverse allocation order  
 → very easy memory management (stack grows upwards)

fixed segment starting point

store top of latest allocation in **stack pointer** (SP)  
 (initialized to starting point)

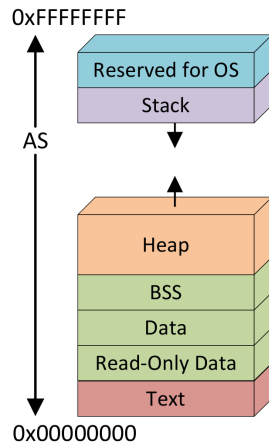
*allocate a* byte data structure: **SP += a; return (SP - a)**

*free a* byte data structure: **SP -= a**

## SEGMENTS — HEAP (DYNAMIC MEMORY ALLOCATION)

some data „randomly“ allocated/free'd  
 two-tier memory allocation:

1. allocate large memory chunk (**heap segment**) from OS  
 - base address + **break pointer** (BRK)  
 - process can get more/give back memory from/to OS
2. dynamically partition chunk into smaller allocations  
 - **malloc/free** can be used in random order  
 - purely user-space, no need to contact kernel



#### Summary

recipe vs. cooking is like program vs. process  
 processes = resource container for OS  
 process feels alone: has own CPU + memory  
 OS implements multiprogramming through rapid process switching

## Process API

### EXECUTION MODEL — ASSEMBLER (SIMPLIFIED)

OS interacts directly with compiled programs

- switch between processes/threads ~> **save/restore** state
- deal with/pass on **signals/exceptions**
- receive **requests** from applications

#### Instructions:


- **mov**: Copy referenced data from second operand to first operand
- **add/sub/mul/div**: Add, ..., from second operand to first operand
- **inc/dec**: increment/decrement register/memory location
- **shl/shr**: shift first operand left/right by amount given by second operand
- **and/or/xor**: calculate bitwise and, ..., of two operands storing result in first
- **not**: bitwise negate operand

### EXECUTION MODEL — STACK (X86)

**stack pointer (SP)**: holds address of stack top (growing downwards)

**stack frames**: larger stack chunks

**base pointer (BP)**: used to organize stack frames



### EXECUTION MODEL — JUMP/BRANCH/CALL COMMANDS (X86)

**jmp**: continue execution at operand address  
**j\$condition**: jump depending on PSW content  
 true ~> jump  
 false ~> continue  
 examples: **je** (jump equal), **jz** (jump zero)  
**call**: push function to stack and jump to it  
**return**: return from function (jump to return address)

### EXECUTION MODEL — APPLICATION BINARY INTERFACE (ABI)

standardizes binary interface between programs, modules, OS:

- executable/object file layout
- calling conventions
- alignment rules

**calling conventions**: standardize exact way function calls are implemented  
 ~> interoperability between compilers

### EXECUTION MODEL — CALLING CONVENTIONS (X86)

function call (caller):

1. save local scope state
2. set up parameters where function can find them
3. transfer control flow

function call (called function):

1. set up new local scope (local variables)
2. perform duty
3. put return value where caller can find it
4. jump back to caller (IP)

### PASSING PARAMETERS TO THE SYSTEM

parameters are passed through **system calls**

call number + specific parameters must be passed

parameters can be transferred through

- **CPU registers** (~6)
- **Main Memory** (heap/stack - more parameters, data types)

ABI specifies how to pass parameters

**return code** needs to be returned to application

- **negative**: error code
- **positive + 0**: success
- usually returned via A+D registers

### SYSTEM CALL HANDLER

implements the actual service called through a syscall:

1. saves tainted registers
2. reads passed parameters
3. sanitizes/checks parameters
4. checks if caller has enough permissions to perform the requested action
5. performs requested action in behalf of the caller
6. returns to caller with success/error code

### PROCESS API — CREATION

process creation events:

1. system initialization
2. process creation syscall
3. user requests process creation
4. batch job-initiation

events map to two mechanisms:

1. Kernel spawns initial user space process on boot (Linux: **init**)
2. User space processes can spawn other processes (within their quota)



## PROCESS API — CREATION (POSIX)

**PID:** identifies process

**pid = fork():** duplicates current process:

- returns 0 to new child
- returns new **PID** to parent
- ~> child and parent independent after **fork**

**exec(name):** replaces own memory based on executable file

- **name** specifies binary executable file

**exit(status):** terminates process, returns **status**

**pid = waitpid(pid, &status):** wait for child termination

- **pid:** process to wait for
- **status:** points to data structure that returns information about the process (e.g., exit status)
- passed **pid** is returned on success, -1 otherwise

**process tree:** processes create child processes, which create child processes, ...

- parent and child execute concurrently
- parent waits for child to terminate (collecting the exit state)

## DAEMONS


= program designed to run in the background

detached from parent process after creation, reattached to process tree root (**init**)

## PROCESS STATES

**blocking:** process does nothing but wait

- usually happens on syscalls (OS doesn't run process until event happens)



## PROCESS TERMINATION

different termination events:

1. normal exit (voluntary)
  - **return** 0 at end of **main**
  - **exit**(0)
2. error exit (voluntary)
  - **return** **x** (**x** ≠ 0) at end of **main**
  - **exit**(**x**) (**x** ≠ 0)
  - **abort**()
3. fatal error (involuntary)
  - OS kills process after exception
  - process exceeds allowed resources
4. killed by another process (involuntary)
  - another process sends kill signal (only as parent process or administrator)

## EXIT STATUS

voluntary exit: process returns exit status (integer)

resources not completely free'd after process terminates

~> **Zombie** or **process stub** (contains exit status until collected via **waitpid**)

**Orphans:** Processes without parents

- usually adopted by **init**
- some systems kill all children when parent is killed

exit status on involuntary exit:

- Bits 0-6: signal number that killed process (0 on normal exit)
- Bit 7: set if process was killed by signal
- Bits 8-15: 0 if killed by signal (exit status on normal exit)

# Threads

## PROCESSES VS. THREADS

**Traditional OS:** each process has

- own address space
- own set of allocated resources
- *one* thread of execution (= one execution state)

**Modern OS:** processes + threads of execution handled more flexibly

- *processes* provide abstraction of address space and resources
- *threads* provide abstraction of execution states of that address space

**Exceptions:**

- sometimes different threads have different address spaces
- *Linux:* threads = regular processes with shared resources and AS regions

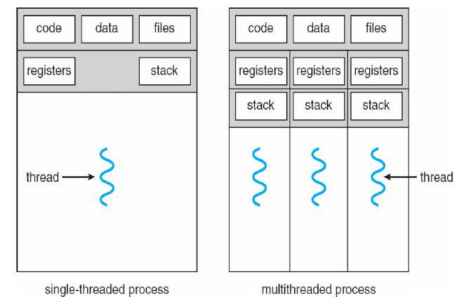
## THREADS — WHY?

many programs do multiple things at once (e.g. web server)

~> writing program as many sequential threads may be easier than with blocking operations

**Processes:** rarely share data (if, then explicitly)

**Threads:** closely related, share data



## THREADS — POSIX

**PThread:** base object with

- *identifier* (thread ID, TID)
- *register set* (including IP and SP)
- *stack area* to hold execution state

**Pthread\_create:** create new thread

- Pass: *pointer* to **pthread\_t** (will hold TID after successful call)
- Pass: *attributes, start function, arguments*
- Returns: 0 on success, error value else

**Pthread\_exit:** terminate calling thread

- Pass: exit code (casted to void pointer)
- Free's resources (e.g. stack)

**Pthread\_join:** wait for specified thread to exit

- Pass: **pthread\_t** to wait for (or -1 for any thread)
- Pass: pointer to pointer for exit code
- Returns: 0 on success, error value else

**Pthread\_yield:** release CPU to let another thread run

## THREADS — PROBLEMS

**Processes vs. Threads:**

- *Processes:* only share resources explicitly
- *Threads:* more shared state → more can go wrong

**Challenges:** programmer needs to take care of

- *activities:* dividing, ordering, balancing
- *data:* dividing
- *shared data:* access synchronizing

## PCB vs. TCB

**PCB** (*process control block*): information needed to implement processes

- always known to OS

**TCB** (*thread control block*): per thread data

- OS knowledge depends on *thread model*

| PCB              | TCB                 |
|------------------|---------------------|
| Address space    | Instruction pointer |
| Global variables | Registers           |
| Open files       | Stack               |
| Child processes  | State               |
| Pending alarms   |                     |

## THREAD MODELS

**Kernel Thread**: known to OS kernel

**User Thread**: known to process

**N:1-Model**: kernel only knows one of possibly multiple threads

- N:1 user threads = *user level threads* (ULT)

**1:1-Model**: each user thread maps to one kernel thread

- 1:1 user threads = *kernel level threads* (KLT)

**M:N-Model** (hybrid model): flexible mapping of user threads to less kernel threads

### THREAD MODELS — N:1

Kernel only manages process → multiple threads unknown to kernel

Threads managed in user-space library (e.g. GNU Portable Threads)

**Pro:**

- + faster thread management operations (up to 100 times)
- + flexible scheduling policy
- + few system resources
- + useable even if OS doesn't support threads

**Con:**

- no parallel execution
- whole process blocks if one user thread blocks
- reimplementing OS parts (e.g. scheduler)


**Stack:**

- main stack known to OS used by thread library
- own execution state (= stack) dynamically allocated by user thread library for each thread
- possibly own stack for each exception handler

**Heap:**

- concurrent heap use possible
- *Attention*: not all heaps are reentrant

**Data**: divided into BSS, data and read-only data here as well



### THREAD MODELS — 1:1

kernel knows + manages every thread

**Pros:**

- + real parallelism possible
- + threads block individually

**Cons:**

- OS manages every thread in system (TCB, stacks,...)
- Syscalls needed for thread management
- scheduling fixed in OS


**Stack:**

- own execution state (= stack) for every thread
- possibly own stack for (each) exception handler

**Heap:**

- parallel heap use possible
- *Attention*: not all heaps are thread-safe
- if thread-safe: not all heap implementations perform well with many threads

**Data**: divided into BSS, data and read-only data here as well



### THREAD MODELS — M:N

**Principle**:  $M$  ULTs are mapped to (at most)  $N$  KLT

- *Goal*: pros of ULT and KLT — non-blocking with quick management
- create sufficient number of KLTs and flexibly allocate ULTs to them
- *Idea*: if ULT blocks ULTs can be switched in userspace

**Pros:**

- + flexible scheduling policy
- + efficient execution

**Cons:**

- hard to debug
- hard to implement (e.g. blocking, number of KLTs,...)

**Implementation — Upcalls:**

- kernel notices that thread will block → sends signal to process
- upcall notifies process of thread id and event that happened
- exception handler of process schedules a different process thread
- kernel later informs process that blocking event finished via other upcall

## Scheduling

### MOTIVATION

$K$  jobs ready to run,  $K > N \geq 1$  CPUs available

**Scheduling Problem:**

- Which jobs should kernel assign to which CPUs?
- When should it make decision?

### DISPATCHER

**Dispatcher**: performs actual process switch

- mechanism
- save/restore process context
- switch to user mode

**Scheduler**: selects next process to run based on *policy*

### VOLUNTARY YIELDING VS. PREEMPTION

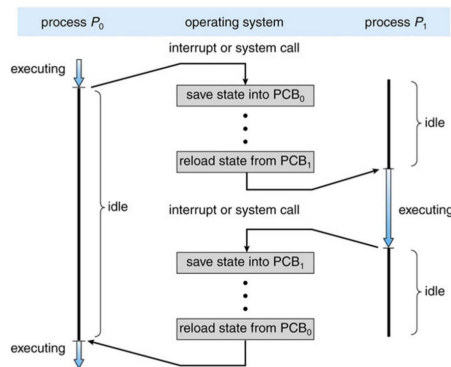
kernel responsible for CPU switch

kernel doesn't always run → can only dispatch different process when invoked

**cooperative multitasking**: running process performs *yield* syscall → kernel switches process

**preemptive scheduling:**

- kernel invoked in certain time intervals
- kernel makes scheduling decisions after every time-slice

**SCHEDULING — PROCESS STATES**

**new:** process was created but did not run yet

**running:** instructions are currently being executed

**waiting:** process is waiting for some event

**ready:** process is waiting to be assigned a processor

**terminated:** process has finished execution

**SCHEDULING — LONG-TERM VS. SHORT-TERM**

**Short-term scheduler** (CPU Scheduler, focused on in this lecture):

- selects process to run next, allocates CPU
- invoked frequently (ms)  $\leadsto$  must be fast

**Long-term scheduler** (job scheduler):

- selects process to be brought into ready queue
- invoked very infrequently (s, m)  $\leadsto$  can be slow
- controls degree of multiprogramming

**SCHEDULING QUEUES**

**job queue:** set of all processes in system

**ready queue:** process in main memory, ready or waiting

**device queue:** processes waiting for I/O device

**SCHEDULING POLICIES — CATEGORIES****batch scheduling:**

- still widespread in business (payroll, inventory,...)
- no users waiting for quick response
- non-preemptive algorithms acceptable  $\rightarrow$  less switches  $\rightarrow$  less overhead

**interactive scheduling:**

- need to optimize for response time
- preemption essential to keep processes from hogging CPU

**real-time scheduling:**

- guarantee job completion within time constraints
- need to be able to plan when which process runs + how long
- preemption not always needed

**SCHEDULING POLICIES — GOALS****General:**

- **fairness:** give each process fair share of CPU
- **balance:** keep all parts of system busy

**batch scheduling:**

- **throughput:** number of processes that complete per time unit
- **turnaround time:** time from job submission to job completion
- **CPU utilization:** keep CPU as busy as possible

**interactive scheduling:**

- **waiting time:** reduce time a process waits in waiting queue
- **response time:** time from request to first response

**real-time scheduling:**

- **meeting deadlines:** finishing jobs in time
- **predictability:** minimize jitter

**SCHEDULING POLICIES — FIRST COME FIRST SERVED**

intuitively clear

**Example:** 3 processes arrive at time 0 in the order  $P_1, P_2, P_3$

| Process | Burst time | Turnaround time |
|---------|------------|-----------------|
| $P_1$   | 24         | 24              |
| $P_2$   | 3          | 27              |
| $P_3$   | 3          | 30              |

$\leadsto$  average turnaround time 27  $\rightarrow$  can we do better?

**Conclusion:** if processes would arrive in order  $P_2, P_3, P_1$ , average turnaround time would be 13

$\leadsto$  good scheduling can reduce turnaround time

**SCHEDULING POLICIES — SHORTEST JOB FIRST**

**Benefits:** optimal average turnaround/waiting/response time

**Challenge:** cannot know job lengths in advance

**Solution:** predict length of next CPU burst for each process

$\leadsto$  schedule process with shortest burst next

**Burst Estimation:** exponential averaging

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n$$

( $t_n$ : actual length of  $n$ -th CPU burst,  $\tau_{n+1}$ : predicted length of next CPU burst,  $0 \leq \alpha \leq 1$ )

**PROCESS BEHAVIOUR — CPU BURSTS**

CPU bursts exist because processes wait for I/O

**CPU-bound processes:** spends more time doing computations

$\leadsto$  few very long CPU bursts

**I/O-bound processes:** spends more time doing I/O

$\leadsto$  many short CPU bursts

**SCHEDULING POLICIES — PREEMPTIVE SHORTEST-JOB-FIRST**

SJF optimizes waiting/response time

$\leadsto$  what about throughput?

**Problem:** CPU-bound jobs hold CPU until exit or I/O  $\rightarrow$  poor I/O utilization

**Idea:** SJF, but preempt periodically to make new scheduling decision

- each time slice: schedule job with shortest remaining time next
- alternatively: schedule job with shortest next CPU burst

**SCHEDULING POLICIES — ROUND ROBIN**

**Problem:** batch schedulers suffer from starvation and don't provide fairness

**Idea:** each process runs for small CPU time unit

- **time quantum/time slice** length: usually 10-100ms
- preempt processes that have not blocked by end of time slice
- append current thread to end of run queue, run next thread

**Caution:** time slice length needs to balance interactivity and overhead!

$\rightarrow$  if time slice length in the area of dispatch time, 50% of CPU time wasted for process switching

**SCHEDULING POLICIES — VIRTUAL ROUND ROBIN**

**Problem:** RR is unfair for I/O-bound jobs: they block before using up time quantum

**Idea:** put jobs that didn't use up their quantum in additional queue

- store share of unused time-slice
- give those jobs additional queue priority
- put them back into normal queue afterwards

### SCHEDULING POLICIES — (STRICT) PRIORITY SCHEDULING

**Problem:** not all jobs are equally important

~> different priorities (e.g., 4)

**Solution:** associate priority number with each process

- RR for each priority
- *aging*: old low priority processes get executed before new higher priority processes

### SCHEDULING POLICIES — MULTI-LEVEL FEEDBACK QUEUE

**Problem:** context switching expensive

~> trade-off between interactivity and overhead?

**Goals:**

- higher priority for I/O jobs (usually don't use up quantum)
- low priority for CPU jobs (rather run them longer)

**Idea:** different queues with different priorities and time slice lengths

- schedule queues with (static) priority scheduling
- double time slice length in each next-lower priority
- process to higher priority when they don't use up quantum repetitively
- process to lower priority when they use up quantum repetitively

### SCHEDULING PRINCIPLES — PRIORITY DONATION

**Problem:** Process B (higher priority) waits for process A (lower priority)

~> B has now effectively lower priority

**Solution:** *priority donation*

- give A priority of B as long as B waits for A
- if C, D, E wait for B → A gets highest priority of B, C, D, E

### SCHEDULING POLICIES — LOTTERY SCHEDULING

issue number of lottery tickets to processes (amount depending on priority)

amount of tickets controls average proportion of CPU for each process

**Scheduling:** scheduler draws random number  $N$ , process with  $N$ -th ticket is executed

processes can transfer tickets to other processes if they wait for them