

Page Faults

PAGE FAULTS — HANDLING

Cause: access to page currently not present in main memory
→ exception, invoking OS

Process:

- OS checks access validity (requiring additional info)
- get empty frame
- load contents of requested page from disk into frame
- adapt page table
- set present bit of respective entry
- restart instruction causing page fault

PAGE FAULTS — LATENCY

fault rate $0 \leq p \leq 1$

- $p = 0$: no page faults
- $p = 1$: every reference leads to page fault

effective access time (EAT):

$$\text{EAT} = (1 - p) * \text{memory access} + p * (\text{PF overhead} + \text{PF service time} + \text{restart overhead})$$

PAGE FAULTS — PERFORMANCE IMPACT

memory access time: 200ns

average page fault service time: 8ms

→ 1:1000 access-page-fault-rate → $\text{EAT} = 8.2\mu\text{s} \Rightarrow \text{slowdown by factor 40!}$

PAGE FAULTS — CHALLENGES

what to eject?

- how to allocate frames among processes?
- which particular process's pages to keep in memory?
- see *page frame allocation*

what to fetch?

- what if block size \neq page size?
- just one page needed? prefetch more?

process resumption?

- need to save state + resume
- process might have been in middle of instruction

PAGE FAULTS — WHAT TO FETCH?

bring in page causing fault

pre-fetch surrounding pages?

- reading two disk blocks is approximately as fast as reading one
- as long as there is no track/head switch, seek (disk) time dominates
- application exhibits spatial locality = big win

pre-zero pages?

- don't want to leak information between processes
- need 0-filled pages for stack, heap, .bss, ...
- *zero on demand?*
- keep pool of 0-pages filled in background when CPU is idle?

PAGE FAULTS — PROCESS RESUMPTION?

hardware provides info about page fault

(intel: `%cr2` contains faulting virtual address)

context: OS needs to figure out fault context:

- read or write?
- instruction fetch?
- user access to kernel memory?

idempotent instructions: easy:

- re-do load/store instructions
- re-execute instructions accessing only one address

complex instructions: must be re-started

- some CISC instructions are hard to restart (e.g., block move of overlapping areas)
- *solutions:*
 - touch relevant pages before operation starts
 - keep modified data in registers → page faults can't take place
 - design ISA such that complex operations can execute partially → consistent page fault state