

Page Replacement Policies

Page Replacement — Naive

- **step 1: save/clear victim page:**
 - drop page if fetched from disk and clean
 - *dirty*: write back modifications if from disk and dirty (unless `MAP_COPY`)
 - *non-dirty*: write page file/swap partition otherwise (e.g., stack, heap memory)
- **step 2: unmap page from old AS:** invalidate PTE, flush cache
- **step 3: prepare new page:** null page or load new contents
- **step 4: map page frame into new AS:** invalidate PTE, flush cache

Page Replacement — Buffering

- **Problem:** naive page replacement encompasses two I/O transfers
→ both operations block page fault from completing
- **Goal:** reduce I/O from critical page fault path to speed up page faults
- **Idea:** keep pool of free page frames (*pre-cleaning*):
 - *on page fault*: use page frame from free pool
 - *cleaning*: daemon cleans, reclaims and scrubs pages for free pool in background→ smooths out I/O, speeds up paging significantly
- **Remaining problem:** which pages to select as victims?
 - *goal*: identify page that has left working set of its processes, add to free pool
 - *success metric*: low overall page fault rate

Page Replacement — FIFO

- **Idea:** evict oldest fetched page in system
- **Belady's Anomaly:** using FIFO, for every number n of page frames you can construct a reference string that performs worse with $n + 1$ frames
→ with FIFO it is possible to get more page faults with more page frames!

Page Replacement — oracle

- = optimal replacement strategy: replace page with next reference furthest in future
- **Problem:** future unpredictable
 - **However:** good metric to check how well other algorithms perform

Page Replacement — LRU

- **Goal:** approximate oracle page replacement
- **Idea:** past often predicts future well
- **Assumption:** page used furthest in past is used furthest in future
- **Cycle counter implementation:**
 - have MMU write CPU's time stamp counter to PTE on every access
 - *page fault*: scan all PTEs to find oldest counter value
 - *advantage*: cheap at access if done in HW
 - *disadvantage*: memory traffic for scanning
- **Stack implementation:**
 - keep doubly linked list of all page frames
 - move each referenced page to tail of list
 - *advantage*: can find replacement victim in $O(1)$
 - *disadvantage*: need to change 6 pointers at every access
- **~ No silver bullet:**
 - *observation*: predicting future based on past is not precise
 - *conclusion*: relax requirements — maybe perfect LRU isn't needed? ⇒ approximate LRU

LRU Approximation — Clock page replacement

- aka *second chance page replacement*
- **Precondition:** MMU sets reference bit in PTE
 - supported natively by most hardware
 - can easily emulate in systems with software managed TLB (e.g., MIPS)
- **Store:** keep all pages in circular FIFO list
- **Searching** for victim: scan pages in FIFO's order
 - if reference bit = 0 → use page as victim and advance
 - if reference bit = 1 → set to 0, continue scanning
- **Problem:** large memory → most pages referenced before scanned
 - *solution*: use 2 arms, leading arm clears reference bit, trailing arm selects victim

Replacement Strategies — other

- **Random Eviction:** pick random victim
 - dirt simple
 - not overly horrible in reality
- **Larger counter:** use n -bit reference counter instead of reference bit

- *least frequently used*: rarely used page not in a working set → replace page with smallest count
- *most frequently used*: page with smallest count probably just brought in → replace page with largest count
- neither LFU nor MDU are common (no such hardware, not that great)

Summary

- victim page frame needs to be selected by OS when handling page faults
 - evicting page frame after page fault happens = not a good idea
 - page buffering keeps eviction out of critical path
- different victim selection policies exist
 - FIFO → Belady's Anomaly
 - Oracle → cannot predict the future
 - Random → unpredictable, never great but rarely very bad
 - LRU → hard to implement efficiently