

Memory Allocation

Memory Allocation — Dynamic

- = allocate + free memory chunks of arbitrary size at arbitrary points in time
 - almost every program uses it (heap)
 - don't have to statically specify complex data structures
 - can have data grow as function of input size
 - kernel itself uses dynamic memory allocation for its data structures
- **Implementation:** has huge impact on performance, both in user and kernel space
- **Fact:** it is impossible to construct memory allocator that always performs well
 - need to understand trade-offs to pick good allocation strategy

Dynamic Memory Allocation — Principle

- **Initial:** pool of free memory
- **Tasks:**
 - satisfy arbitrary **allocate** + **free** requests from pool
 - track which parts are in use/are free
- **Restrictions:**
 - cannot control order/number of requests
 - cannot move allocated regions → fragmentation = core problem!

Dynamic Memory Allocation — Bitmap

- **Idea:**
 - divide memory in allocation units of fixed size
 - use bitmap to keep track if allocated (1) or free (0)
- **Problem:** needs additional data structure to store allocation length (otherwise cannot infer whether two adjacent allocations belong together or not from bitmap)

Dynamic Memory Allocation — List

- **Method 1:** use one list-node for each allocated data
 - *extra space* needed for list
 - allocation lengths already stored
- **Method 2:** use one list-node for each unallocated data
 - can keep list in unallocated area (store size of free area + pointer to next free area in free area)
 - *additional data structure* needed to store allocation lengths
 - can search for free space with low overhead
- **Method 3:** both

Dynamic Memory Allocation — Problems

- **Fragmentation** is hard to handle
- **Factors** needed for fragmentation to occur:
 - *different lifetimes*
 - *different sizes*
 - *inability to relocate previous allocations*
- all fragmentation factors present in dynamic memory allocators!

Allocation — Best fit vs. Worst fit

- **Idea:** keep large free memory chunks together for larger allocation requests that may arrive later
- **Best-fit:** allocate smallest free block large enough to store allocation request
 - must search entire list
 - *problem:* sawdust — remainder so small that over time left with unusable sawdust everywhere
 - *idea:* minimize sawdust by turning strategy around
- **Worst-fit:** allocate largest free block
 - must search entire list
 - *reality:* worse fragmentation than best-fit

Allocation — First fit

- **Idea:** if fragmentation occurs with best and worst fit, optimize for allocation speed
- **Principle:** allocate first hole big enough
 - fastest allocation policy
 - produced leftover holes of variable size
 - *reality:* almost as good as best-fit

First Fit — Variants

- **first-fit sorted by address order**
- **LIFO first-fit**

- **next fit**

Allocation — Buddy allocator

- **Idea:** allocate memory in powers of 2
 - all chunks have fixed 2^n -size → allocation request rounded up to next-higher power of 2
 - all chunks naturally aligned
- **no sufficiently small block available:**
 - select larger available chunk, split into two same-sized buddies
 - continue until appropriately sized chunk is available
- **two buddies both free** (2^n): merge to 2^{n+1} -chunk

Real Program Patterns

- **Ramps:** accumulate data monotonically over time
- **Peaks:** allocate many objects, use briefly, then free all
- **Plateaus:** allocate many objects, use for long time

Allocation — Slabs

- kernel often allocates/frees memory for few, specific data objects of fixed size
- **Slab:** multiple pages of contiguous physical memory
 - linux: uses buddy allocator as underlying allocator for slabs
- **Cache:** one or multiple slabs
 - stores only one kind of object (fixed size)

Summary

- dynamic memory means allocating and freeing memory chunks of different sizes at any time
- impossible to construct memory allocator that always performs well
- typical dynamic memory data structures:
 - bitmaps
 - lists
- simple, well-performing allocation strategies:
 - best-fit
 - first-fit
- advanced strategies:
 - buddy-allocator
 - slab-allocator