

Page Faults

PAGE FAULTS — HANDLING

Cause: access to page currently not present in main memory
→ exception, invoking OS

Process:

- OS checks access validity (requiring additional info)
- get empty frame
- load contents of requested page from disk into frame
- adapt page table
- set present bit of respective entry
- restart instruction causing page fault

PAGE FAULTS — LATENCY

fault rate $0 \leq p \leq 1$

- $p = 0$: no page faults
- $p = 1$: every reference leads to page fault

effective access time (EAT):

$$\text{EAT} = (1 - p) * \text{memory access} + p * (\text{PF overhead} + \text{PF service time} + \text{restart overhead})$$

PAGE FAULTS — PERFORMANCE IMPACT

memory access time: 200ns

average page fault service time: 8ms

→ 1:1000 access-page-fault-rate → $\text{EAT} = 8.2\mu\text{s} \Rightarrow \text{slowdown by factor 40!}$

PAGE FAULTS — CHALLENGES

what to eject?

- how to allocate frames among processes?
- which particular process's pages to keep in memory?
- see *page frame allocation*

what to fetch?

- what if block size \neq page size?
- just one page needed? prefetch more?

process resumption?

- need to save state + resume
- process might have been in middle of instruction

PAGE FAULTS — WHAT TO FETCH?

bring in page causing fault

pre-fetch surrounding pages?

- reading two disk blocks is approximately as fast as reading one
- as long as there is no track/head switch, seek (disk) time dominates
- application exhibits spatial locality = big win

pre-zero pages?

- don't want to leak information between processes
- need 0-filled pages for stack, heap, .bss, ...
- *zero on demand?*
- keep pool of 0-pages filled in background when CPU is idle?

PAGE FAULTS — PROCESS RESUMPTION?

hardware provides info about page fault

(intel: `%cr2` contains faulting virtual address)

context: OS needs to figure out fault context:

- read or write?
- instruction fetch?
- user access to kernel memory?

idempotent instructions: easy:

- re-do load/store instructions
- re-execute instructions accessing only one address

complex instructions: must be re-started

- some CISC instructions are hard to restart (e.g., block move of overlapping areas)
- *solutions:*
 - touch relevant pages before operation starts
 - keep modified data in registers → page faults can't take place
 - design ISA such that complex operations can execute partially → consistent page fault state

MEMORY-MAPPED FILES — OTHER ISSUES

I/O mapping: mapping disk block to page in memory allows file I/O to be treated as routing memory

- *initial:* read page-sized portion of file from file system to physical page
- *subsequent read/write:* treated as ordinary memory access
- *simplifies* file access, file I/O through memory instead of syscalls
- *memory-file sharing:* several processes can map to same file

SHARED DATA SEGMENTS

implementation:

- temporary, asynchronous memory-mapped files
- shared pages (with allocated space on backing store)

copy on write (COW):

- allows both parent and child process to initially share same memory pages
- only modified pages are copied → more efficient process creation

PAGE FRAME ALLOCATION — LOCAL VS. GLOBAL

global: all frames considered for replacement

- does not consider page ownership
- one process cannot get another process's frame
- does not protect process from a process that hogs all memory

local: only frames of faulting process are considered for replacement

- isolates processes/users
- separately determine how many frames each process gets

FIXED ALLOCATION — EQUAL VS. PROPORTIONAL

equal: all processes get same amount of frames

proportional: allocate according to process size

$$s_i := \text{size of process } p_i, S := \sum s_i, m := \text{total number of frames}$$

$$\Rightarrow a_i := \frac{s_i}{S} m \text{ allocation for } p_i$$

FIXED ALLOCATION — PRIORITY ALLOCATION

= proportional allocation scheme using priorities rather than size

on page fault of P_i :

- select one of its frames for replacement or
- select frame from process with lower priority

MEMORY LOCALITY

problem: background storage much slower than memory

- paging extends memory size using background storage
- *goal:* run near memory speed, not near background storage speed

Pareto principle: applies to working sets of processes

- 10% of memory gets 90% of references
- *goal:* keep those 10% in memory, rest on disk
- *problem:* how to identify those 10%?

THRASHING

problem: system is busy swapping pages in and out

- each time one page is brought in, another page, whose contents will soon matter, is thrown out
- *effect:* low CPU utilization, processes wait for pages to be fetched from disk
- *consequence:* OS thinks that it needs higher degree of multiprogramming

reasons:

- *no temporal locality* of access pattern — process doesn't follow Pareto principle
- *too much multiprogramming:* each process fits individually, but too many for system
- *memory too small* to hold hot memory of a single process (the 10%)
- *bad page replacement policy*

WORKING-SET MODEL

Δ := working-set window (fixed number of page references; e.g., 10000 instructions)

WSS_i := working set of process P_i

- total number of pages referenced in most recent Δ (varies in time)
- Δ
 - too small \Rightarrow will not encompass entire locality
 - too large \Rightarrow will encompass several localities
 - $= \infty \Rightarrow$ will encompass entire program

$D := \sum WSS_i$ = total demand for frames

- $D > m \rightsquigarrow$ **thrashing**
- $D > m \Rightarrow$ suspend a process

WORKING SET — KEEPING TRACK

perfect: replace page that is referenced furthest in the future (*oracle*)

idea: predict future from past

- record page references from past and extrapolate into future
- *problem:* too expensive to make ordered list of all page references at runtime

idea: sacrifice precision for speed

- MMU sets *reference bit* in respective page table entry every time a page is referenced
- set timer to scan all page table entries for reference bits

PAGE FAULT FREQUENCY — ALLOCATION SCHEME

goal: establish acceptable page fault rate

- *actual rate too low* \rightarrow give frames to other process
- *actual rate too high* \rightarrow allocate more frames to process

PAGE FETCH POLICY — DEMAND PAGING

idea: only transfer pages raising page faults

advantages:

- only transfer what is needed
- less memory needed by process \rightarrow higher multiprogramming degree possible

disadvantages:

- many initial page faults when task starts
- more I/O operations \rightarrow more I/O overhead

PAGE FETCH POLICY — PRE-PAGING

idea: speculatively transfer pages to RAM

- at every page fault: speculate what else should be loaded
- e.g., load entire text section when process starts

advantage: improves disk I/O throughput

disadvantages:

- wastes I/O bandwidth if page is never used
- can destroy working set of other processes in case of page stealing

Summary

paging simulates a memory size of the size of the virtual memory

when pages are filled via page faults, OS needs to answer some questions:

- what to eject?
- what to fetch?
- how to resume process?

different strategies to allocate frames and replace pages:

- local vs. global allocation
- fixed vs. proportional vs. priority allocation

thrashing must be prevented by taking working sets of active processes into account