

Introduction to Operating Systems

What's an OS?

The OS is a layer between applications and hardware to ease development.

- **Abstraction.** provides abstraction for applications:
 - manages + hides hardware details
 - uses low-level interfaces (not available to applications)
 - multiplexes hardware to multiple programs (*virtualization*)
 - makes hardware use efficient for applications
- **Protection.**
 - from processes using up all resources (*accounting, allocation*)
 - from processes writing into other processes memory
- **Resource Management.**
 - manages + multiplexes hardware resources
 - decides between conflicting requests for resource use
 - *goal*: efficient + fair resource use
- **Control.**
 - controls program execution
 - prevents errors and improper computer use

→ no universally accepted definition

Hardware Overview

- **Bus:** CPU(s)/devices/memory (conceptually) connected to common bus
 - CPU(s)/devices competing for memory cycles/bus
 - all entities run concurrently
 - *today*: multiple buses
- **Device controller:** has local buffer and is in charge of particular device
- **Interplay:**
 1. CPU issues commands, moves data to devices
 2. Device controller informs APIC that it has finished operation
 3. APIC signals CPU
 4. CPU receives device/interrupt number from APIC, executes handler

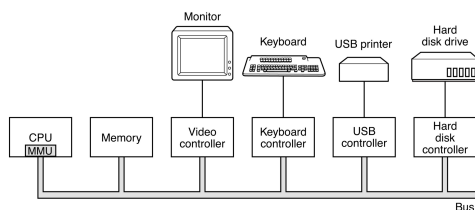


Figure 1: Traditional bus design.

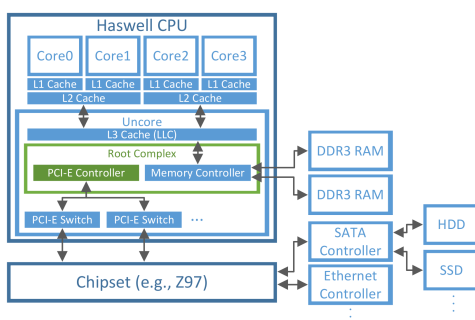


Figure 2: Modern bus design.

Central Processing Unit (CPU) — Operation

- **Principle:**
 1. *fetches* instructions from memory,
 2. *executes* them
- **During execution:** (meta-)data is stored in CPU-internal registers, i.e.
 - general purpose registers
 - floating point registers
 - instruction pointer (IP)
 - stack pointer (SP)
 - program status word (PSW)

CPU — Modes of Execution

- **User mode** (x86: Ring 3/CPL 3):
 - only non-privileged instructions may be executed
 - cannot manage hardware → *protection*
- **Kernel mode** (x86: Ring 0/CPL 0):
 - all instructions allowed
 - can manage hardware with *privileged instructions*

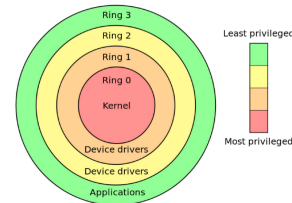


Figure 3: The different protection layers in the ring model.

Random Access Memory (RAM)

- **Principle:** keeps currently executed instructions + data
- **Connectivity:**
 - *today*: CPUs have built-in **memory controller**
 - **CPU caches**: “wired” to CPU
 - **RAM**: connected via pins
 - **PCI-E switches**: connected via pins

Caches

- **Problem:** RAM delivers instructions/data slower than CPU can execute
- **Locality principle:**
 - *spatial locality*: future refs often near previous accesses (e.g. next byte in array)
 - *temporal locality*: future refs often at previously accessed ref (e.g. loop counter)
- **Solution:** *caching* helps mitigating this memory wall
 1. *copy* used information temporarily from slower to faster storage
 2. *check* faster storage first before going down *memory hierarchy*
 3. if *not found*, data is copied to cache and used from there
- **Access latency:**
 - *register*: ~1 CPU cycle
 - *L1 cache* (per core): ~4 CPU cycles
 - *L2 cache* (per core pair): ~12 CPU cycles
 - *L3 cache/LLC* (per uncore): ~28 CPU cycles (~25 GiB/s)
 - *DDR3-12800U RAM*: ~28 CPU cycles + ~50ns (~12 GiB/s)

Device controlling

- **Device controller:** controls device, accepts commands from OS via *device driver*
- **Device registers/memory:**
 - *control* device by writing device registers
 - *read* status of device by reading device registers
 - *pass data* to device by reading/writing device memory
- **Device registers/memory access:**
 1. **port-mapped IO (PMIO):** use special CPU instructions to access port-mapped registers/memory
 2. **memory-mapped IO (MMIO):**
 - use same address space for RAM and device memory
 - some addresses map to RAM, others to different devices
 - access device's memory region to access device registers/memory
 3. **Hybrid:** some devices use hybrid approaches using both

Summary

- The OS is an **abstraction** layer between applications and hardware (multiplexes hardware, hides hardware details, provides protection between processes/users)
- The CPU provides a **separation** of User and Kernel mode (which are required for an OS to provide protection between applications)
- CPU can execute commands faster than memory can deliver instructions/data — **memory hierarchy** mitigates this memory wall, needs to be carefully managed by OS to minimize slowdowns
- device drivers **control** hardware devices through PMIO/MMIO
- Devices can **signal** the CPU (and through the CPU notify the OS) through interrupts

OS Concepts

OS Invocation

- OS Kernel does **not** always run in background!
- Occasions invoking kernel, switching to kernel mode:
 1. **System calls**: User-Mode processes require higher privileges
 2. **Interrupts**: CPU-external device sends signal
 3. **Exceptions**: CPU signals unexpected condition

System Calls — Motivation

- **Problem**: protect processes from one another
- **Idea**: Restrict processes by running them in user-mode
- **~ Problem**: now processes cannot manage hardware,...
 - who can switch between processes?
 - who decides if process may open certain file?
- **~ Idea**: OS provides **services** to apps
 1. app calls system if service is needed (**syscall**)
 2. OS checks if app is allowed to perform action
 3. if app may perform action and hasn't exceeded quota, OS performs action in behalf of app in kernel mode

System Calls — Examples

- `fd = open(file, how, ...)` – open file for read/write/both
- documented e.g. in `man 2 write`
- overview in `man 2 syscalls`

System Calls vs. APIs

- **Syscalls**: interface between apps and OS services, limited number of well-defined entry points to kernel
- **APIs**: often used by programmers to make syscalls (e.g. `printf` library call uses `write` syscall)
- common APIs: Win32, POSIX, C API

System Calls — Implementation

- **Trap Instruction**: single syscall interface (entry point) to kernel
 - switches CPU to kernel mode, enters kernel in same way for all syscalls
 - *system call dispatcher* in kernel then acts as syscall multiplexer
- **Syscall Identification**: number passed to trap instruction
 - *Syscall Table* maps syscall numbers to kernel functions
 - *Dispatcher* decides where to jump based on number and table
 - programs (e.g. `stdlib`) have syscall number compiled in!
 - ~ never reuse old syscall numbers in future kernel versions

Interrupts

- **Devices**: use interrupts to signal predefined conditions to OS
 - *reminder*: device has “interrupt line” to CPU (e.g. device controller informs CPU that operation is finished)
- **Programmable Interrupt Controller**: manages interrupts
 - interrupts can be *masked* (queued, delivered when interrupt unmasked)
 - queue has finite length ~ interrupts can get lost
- **Examples**:
 1. *timer-interrupt*: periodically interrupts processes, switches to kernel ~ can then switch to different processes for fairness
 2. *network interface card* interrupts CPU when packet was received ~ can deliver packet to process and free NIC buffer
- **Interrupt process**:
 1. CPU looks up *interrupt vector* (= table pinned in memory, contains addresses of all service routines)
 2. CPU transfers control to respective *interrupt service routine* in OS that handles interrupt~ interrupt service routine must first save interrupted process's state (instruction pointer, stack pointer, status word)

Exceptions

- **Motivation**: unusual condition → impossible for CPU to continue processing
- **~ Exception** generated within CPU:
 1. CPU interrupts program, gives kernel control
 2. kernel determines reason for exception
 3. if kernel can resolve problem ~ does so, continues *faulting instruction*
 4. kills process if not

- **Difference to Interrupts**: interrupts can happen in any context, exceptions always occur asynchronous and in process context

OS Concepts — Physical Memory

- up to early 60s:
 - programs loaded and run directly in *physical memory*
 - program too large → partitioned manually into *overlays*
 - OS: swaps overlays between disk and memory
 - different jobs could observe/modify each other

OS Concepts — Address Spaces

- **Motivation**: bad programs/people need to be isolated
- **Idea**: give every job the illusion of having all memory to itself
 - every job has own *address space*, can't name addresses of others
 - jobs always and only use virtual addresses

Virtual Memory — Indirect Addressing

- **MMU**: every CPU has built-in *memory management unit* (MMU)
- **Principle**: translates virtual addresses to physical addresses at every load/store
~ address translation protects one program from another
- **Definitions**:
 - *Virtual address*: address in process' address space
 - *Physical address*: address of real memory

Virtual Memory — Memory Protection

- **Kernel-only Virtual Addresses**
 - kernel typically part of all address spaces
 - ensures that apps can't touch kernel memory
- **Read-only virtual addresses**: can be enforced by MMU
 - allows safe sharing of memory between apps
- **Execute Disable**: can be enforced by MMU
 - makes code injection attacks harder

Virtual Memory — Page Faults

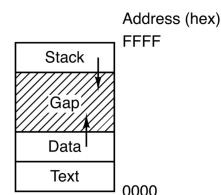
- **Motivation**: not all addresses need to be mapped at all times
 - MMU issues *page fault* exception when accessed virtual address isn't mapped
 - OS handles page faults by loading faulting addresses and then continuing the program
- ~ memory can be *over-committed*: more memory than physically available can be allocated to application
- **Illegal addresses**: page faults also issued by MMU on illegal memory accesses

OS Concepts — Processes

- **Process**: program in execution (“instance” of program)
- each process is associated with
 - **Process Control Block** (PCB): contains information about allocated resources
 - virtual **Address Space** (AS):
 - all (virtual) memory locations a program can name
 - starts at 0 and runs up to a maximum
 - address 123 in AS1 generally ≠ address 123 in AS2
 - indirect addressing ~ different ASes to different programs
- ~ *protection between processes*

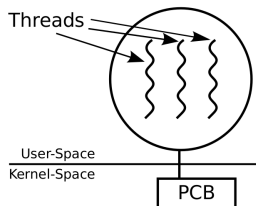
OS Concepts — Address Space Layout

- **Sections**: address spaces typically laid-out in different sections
 - memory addresses between sections *illegal*
 - illegal addresses ~ page fault (*segmentation fault*)
 - OS usually kills process causing segmentation fault
- **Important sections**:
 - *Stack*: function history, local variables
 - *Data*: Constants, static/global variables, strings
 - *Text*: Program code



OS Concepts — Threads

- **Thread:** represents execution state of process (≥ 1 thread per process)
 - *IP:* stores currently executed instruction (address in **text** section)
 - *SP:* stores address of stack top (> 1 threads \rightarrow multiple stacks!)
 - *PSW:* contains flags about execution history (e.g. last calculation was 0 \rightarrow used in following jump instruction)
 - more general purpose registers, floating point registers,...



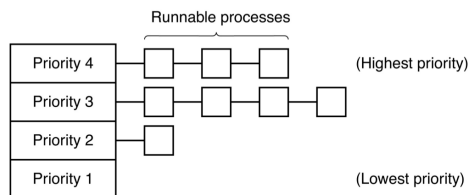
OS Concepts — Policies vs. Mechanisms

- **Mechanism:** implementation of what is done (e.g. commands to write to HDD)
- **Policy:** rules which decide when what is done and how much (e.g. how often, how many resources are used,...)

\rightarrow mechanisms can be reused even when policy changes

OS Concepts — Scheduling

- **Motivation:** multiple processes/threads available \leadsto OS needs to switch between them (for multitasking)
- **Scheduler:** decides which job to run next (*policy*) — tries to
 - provide fairness
 - meet performance goals
 - adhere to priorities
- **Dispatcher:** performs task-switching (*mechanism*)



OS Concepts — Files

- **Motivation:** OS hides peculiarities of file storage, programmer uses device-independent *files/directories*
- **Files:** associate *file name* and *offset* with bytes
- **Directories:** associate *directory names* with directory names or file names
- **File System:** ordered block collection
 - main task: translate (dir name + file name + offset) to block
 - programmer uses file system operations to operate on files (**open, read, seek**)
 - processes can communicate directly through special *named pipe* file (used with same operations as any other file)

OS Concepts — Directory Tree

- **Directories:** form *directory tree/file hierarchy* \rightarrow structure data
- **Root Directory:** topmost directory in tree
- **Path Name:** used to specify file

OS Concepts — Mounting

- **Unix:** common to orchestrate multiple file systems in single file hierarchy
- file systems can be *mounted* on directory
- **Win:** manage multiple directory hierarchies with drive letters (e.g. **C:** \Users)

OS Concepts — Storage Management

- **OS:** provides uniform view of information storage to file systems
 - *Drivers:* hide specific hardware devices \rightarrow hides device peculiarities
 - general interface abstracts physical properties to logical units \rightarrow block
- **Performance:** OS increases I/O performance:
 - *Buffering:* Store data temporarily while transferred
 - *Caching:* Store data parts in faster storage
 - *Spooling:* Overlap one job's output with other job's input

Summary

- **OS:** provides abstractions for and protection between applications
- **Kernel:** does not always run — certain events invoke kernel
 - *syscall:* process asks kernel for service
 - *interrupt:* device sends signal that OS has to handle
 - *exception:* CPU encounters unusual situation
- **Processes:** encapsulate resources needed to run program in OS
 - *threads:* represent different execution states of process
 - *address space:* all memory process can name
 - *resources:* allocated resources, e.g., open files
- **Scheduler** decides which process to run next when multi-tasking
- **Virtual Memory** implements address spaces, provides protection between processes
- **File system** abstracts background store using I/O drivers, provides simple interface (files + directories)

Processes

The Process Abstraction

- **Motivation:** computers (seem to) do "several things at the same time" (quick process switching \rightarrow *multiprogramming*)
- **Model:** *process abstraction* models this concurrency:
 - container contains information about program execution
 - conceptually, every process has own "virtual CPU"
 - execution context is changed on process switch
 - dispatcher switches context when switching processes
 - **context switch:** dispatcher saves current registers/memory mappings, restores those of next process

Process-Cooking Analogy

- Program/Process like Recipe/Cooking
- **Recipe:** lists ingredients, gives algorithm what to do when
 - \leadsto program describes memory layout/CPU instructions
- **Cooking:** activity of using the recipe
 - \leadsto process is activity of executing a program
- multiple similar recipes for same dish
 - \leadsto multiple programs may solve same problem
- recipe can be cooked in different kitchens at the same time
 - \leadsto program can be run on different CPUs at the same time (as different processes)
- multiple people can cook one recipe
 - \leadsto one process can have several worker threads

Concurrency vs. Parallelism

- OS uses concurrency + parallelism to implement multiprogramming
 1. **Concurrency:** multiple processes, one CPU
 - \leadsto not at the same time
 2. **Parallelism:** multiple processes, multiple CPU
 - \leadsto at the same time

Virtual Memory Abstraction — Address Spaces

- every process has own *virtual addresses* (**vaddr**)
- MMU relocates each load/store to *physical memory* (**pmem**)
- processes never see physical memory, can't access it directly
- + MMU can enforce protection (mappings in kernel mode)
- + programs can see more memory than available
 - 80:20 rule: 80% of process memory idle, 20% active
 - can keep working set in RAM, rest on disk
- need special MMU hardware

Address Space (Process View)

- **Motivation:** code/data/state need to be organized within process
 - \leadsto *address space layout*
- **Data types:**
 1. *fixed size* data items
 2. data naturally *freed in reverse allocation order*
 3. data *allocated/freed "randomly"*
- compiler/architecture determine how large int is and what instructions are used in text section (**code**)

- **Loader** determines based on exe file how executed program is placed in memory

Segments — Fixed-Size Data + Code

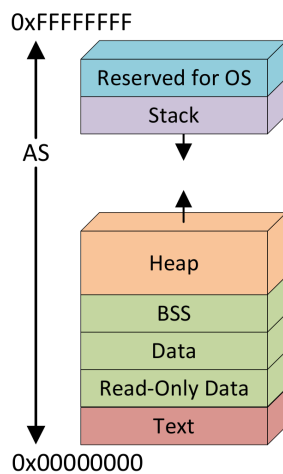
- some data in programs never changes or will be written but never grows/shrinks
→ memory can be statically allocated on process creation
- **BSS segment** (*block started by symbol*):
 - statically allocated variables/non-initialized variables
 - executable file typically contains starting address + size of BSS
 - entire segment initially 0
- **Data segment**: fixed-size, initialized data elements (e.g. global variables)
- **Read-only data segment**: constant numbers, strings
- All three sometimes summarized as one segment
- compiler and OS decide ultimately where to place which data/how many segments exist

Segments — Stack

- some data naturally freed in reverse allocation order
 - very easy memory management (stack grows upwards)
- fixed segment starting point
- store top of latest allocation in **stack pointer** (SP) (initialized to starting point)
- *allocate* **a** byte data structure: `SP += a; return(SP - a)`
- *free* **a** byte data structure: `SP -= a`

Segments — Heap (Dynamic Memory Allocation)

- some data "randomly" allocated/freed
- two-tier memory allocation:
 1. allocate large memory chunk (**heap segment**) from OS
 - base address + **break pointer** (BRK)
 - process can get more/give back memory from/to OS
 2. dynamically partition chunk into smaller allocations
 - `malloc/free` can be used in random order
 - purely user-space, no need to contact kernel



Summary

Processes: recipe vs. cooking = program vs. process

- processes = resource container for OS
- process feels alone (has own CPU and memory)
- OS implements multiprogramming through rapid process switching