

# Introduction to Operating Systems

## What's an OS?

The OS is a layer between applications and hardware to ease development.

- **Abstraction.** provides abstraction for applications:
  - manages + hides hardware details
  - uses low-level interfaces (not available to applications)
  - multiplexes hardware to multiple programs (*virtualization*)
  - makes hardware use efficient for applications
- **Protection.**
  - from processes using up all resources (*accounting, allocation*)
  - from processes writing into other processes memory
- **Resource Management.**
  - manages + multiplexes hardware resources
  - decides between conflicting requests for resource use
  - *goal*: efficient + fair resource use
- **Control.**
  - controls program execution
  - prevents errors and improper computer use

→ no universally accepted definition

## Hardware Overview

- **Bus:** CPU(s)/devices/memory (conceptually) connected to common bus
  - CPU(s)/devices competing for memory cycles/bus
  - all entities run concurrently
  - *today*: multiple buses
- **Device controller:** has local buffer and is in charge of particular device
- **Interplay:**
  1. CPU issues commands, moves data to devices
  2. Device controller informs APIC that it has finished operation
  3. APIC signals CPU
  4. CPU receives device/interrupt number from APIC, executes handler

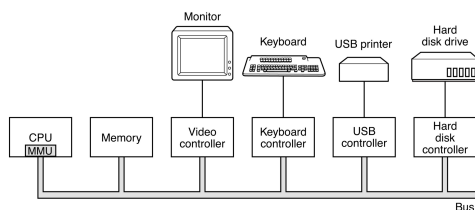


Figure 1: Traditional bus design.

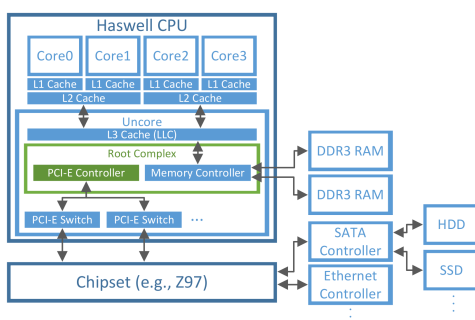


Figure 2: Modern bus design.

## Central Processing Unit (CPU) — Operation

- **Principle:**
  1. *fetches* instructions from memory,
  2. *executes* them
- **During execution:** (meta-)data is stored in CPU-internal registers, i.e.
  - general purpose registers
  - floating point registers
  - instruction pointer (IP)
  - stack pointer (SP)
  - program status word (PSW)

## CPU — Modes of Execution

- **User mode** (x86: Ring 3/CPL 3):
  - only non-privileged instructions may be executed
  - cannot manage hardware → *protection*
- **Kernel mode** (x86: Ring 0/CPL 0):
  - all instructions allowed
  - can manage hardware with *privileged instructions*

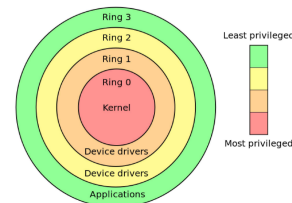


Figure 3: The different protection layers in the ring model.

## Random Access Memory (RAM)

- **Principle:** keeps currently executed instructions + data
- **Connectivity:**
  - *today*: CPUs have built-in **memory controller**
  - **CPU caches**: “wired” to CPU
  - **RAM**: connected via pins
  - **PCI-E switches**: connected via pins

## Caches

- **Problem:** RAM delivers instructions/data slower than CPU can execute
- **Locality principle:**
  - *spatial locality*: future refs often near previous accesses (e.g. next byte in array)
  - *temporal locality*: future refs often at previously accessed ref (e.g. loop counter)
- **Solution:** *caching* helps mitigating this memory wall
  1. *copy* used information temporarily from slower to faster storage
  2. *check* faster storage first before going down *memory hierarchy*
  3. if *not found*, data is copied to cache and used from there
- **Access latency:**
  - *register*: ~1 CPU cycle
  - *L1 cache* (per core): ~4 CPU cycles
  - *L2 cache* (per core pair): ~12 CPU cycles
  - *L3 cache/LLC* (per uncore): ~28 CPU cycles (~25 GiB/s)
  - *DDR3-12800U RAM*: ~28 CPU cycles + ~50ns (~12 GiB/s)

## Device controlling

- **Device controller:** controls device, accepts commands from OS via *device driver*
- **Device registers/memory:**
  - *control* device by writing device registers
  - *read* status of device by reading device registers
  - *pass data* to device by reading/writing device memory
- **Device registers/memory access:**
  1. **port-mapped IO** (PMIO): use special CPU instructions to access port-mapped registers/memory
  2. **memory-mapped IO** (MMIO):
    - use same address space for RAM and device memory
    - some addresses map to RAM, others to different devices
    - access device's memory region to access device registers/memory
  3. **Hybrid**: some devices use hybrid approaches using both

### Summary

- The OS is an **abstraction** layer between applications and hardware (multiplexes hardware, hides hardware details, provides protection between processes/users)
- The CPU provides a **separation** of User and Kernel mode (which are required for an OS to provide protection between applications)
- CPU can execute commands faster than memory can deliver instructions/data — **memory hierarchy** mitigates this memory wall, needs to be carefully managed by OS to minimize slowdowns
- device drivers **control** hardware devices through PMIO/MMIO
- Devices can **signal** the CPU (and through the CPU notify the OS) through interrupts

# OS Concepts

## OS Invocation

- OS Kernel does **not** always run in background!
- Occasions invoking kernel, switching to kernel mode:
  1. **System calls**: User-Mode processes require higher privileges
  2. **Interrupts**: CPU-external device sends signal
  3. **Exceptions**: CPU signals unexpected condition

## System Calls — Motivation

- **Problem**: protect processes from one another
- **Idea**: Restrict processes by running them in user-mode
- **~ Problem**: now processes cannot manage hardware,...
  - who can switch between processes?
  - who decides if process may open certain file?
- **~ Idea**: OS provides **services** to apps
  1. app calls system if service is needed (**syscall**)
  2. OS checks if app is allowed to perform action
  3. if app may perform action and hasn't exceeded quota, OS performs action in behalf of app in kernel mode

## System Calls — Examples

- `fd = open(file, how, ...)` – open file for read/write/both
- documented e.g. in `man 2 write`
- overview in `man 2 syscalls`

## System Calls vs. APIs

- **Syscalls**: interface between apps and OS services, limited number of well-defined entry points to kernel
- **APIs**: often used by programmers to make syscalls (e.g. `printf` library call uses `write` syscall)
- common APIs: Win32, POSIX, C API

## System Calls — Implementation

- **Trap Instruction**: single syscall interface (entry point) to kernel
  - switches CPU to kernel mode, enters kernel in same way for all syscalls
  - *system call dispatcher* in kernel then acts as syscall multiplexer
- **Syscall Identification**: number passed to trap instruction
  - *Syscall Table* maps syscall numbers to kernel functions
  - *Dispatcher* decides where to jump based on number and table
  - programs (e.g. `stdlib`) have syscall number compiled in!
  - ~ never reuse old syscall numbers in future kernel versions

## Interrupts

- **Devices**: use interrupts to signal predefined conditions to OS
  - *reminder*: device has “interrupt line” to CPU (e.g. device controller informs CPU that operation is finished)
- **Programmable Interrupt Controller**: manages interrupts
  - interrupts can be *masked* (queued, delivered when interrupt unmasked)
  - queue has finite length ~ interrupts can get lost
- **Examples**:
  1. *timer-interrupt*: periodically interrupts processes, switches to kernel ~ can then switch to different processes for fairness
  2. *network interface card* interrupts CPU when packet was received ~ can deliver packet to process and free NIC buffer
- **Interrupt process**:
  1. CPU looks up *interrupt vector* (= table pinned in memory, contains addresses of all service routines)
  2. CPU transfers control to respective *interrupt service routine* in OS that handles interrupt~ interrupt service routine must first save interrupted process's state (instruction pointer, stack pointer, status word)

## Exceptions

- **Motivation**: unusual condition → impossible for CPU to continue processing
- **~ Exception** generated within CPU:
  1. CPU interrupts program, gives kernel control
  2. kernel determines reason for exception
  3. if kernel can resolve problem ~ does so, continues *faulting instruction*
  4. kills process if not

- **Difference to Interrupts**: interrupts can happen in any context, exceptions always occur asynchronous and in process context

## OS Concepts — Physical Memory

- up to early 60s:
  - programs loaded and run directly in *physical memory*
  - program too large → partitioned manually into *overlays*
  - OS: swaps overlays between disk and memory
  - different jobs could observe/modify each other

## OS Concepts — Address Spaces

- **Motivation**: bad programs/people need to be isolated
- **Idea**: give every job the illusion of having all memory to itself
  - every job has own *address space*, can't name addresses of others
  - jobs always and only use virtual addresses

## Virtual Memory — Indirect Addressing

- **MMU**: every CPU has built-in *memory management unit* (MMU)
- **Principle**: translates virtual addresses to physical addresses at every load/store  
~ address translation protects one program from another
- **Definitions**:
  - *Virtual address*: address in process' address space
  - *Physical address*: address of real memory

## Virtual Memory — Memory Protection

- **Kernel-only Virtual Addresses**
  - kernel typically part of all address spaces
  - ensures that apps can't touch kernel memory
- **Read-only virtual addresses**: can be enforced by MMU
  - allows safe sharing of memory between apps
- **Execute Disable**: can be enforced by MMU
  - makes code injection attacks harder

## Virtual Memory — Page Faults

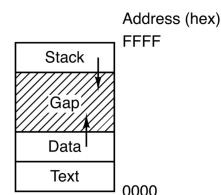
- **Motivation**: not all addresses need to be mapped at all times
  - MMU issues *page fault* exception when accessed virtual address isn't mapped
  - OS handles page faults by loading faulting addresses and then continuing the program
- ~ memory can be *over-committed*: more memory than physically available can be allocated to application
- **Illegal addresses**: page faults also issued by MMU on illegal memory accesses

## OS Concepts — Processes

- **Process**: program in execution (“instance” of program)
- each process is associated with
  - **Process Control Block** (PCB): contains information about allocated resources
  - virtual **Address Space** (AS):
    - all (virtual) memory locations a program can name
    - starts at 0 and runs up to a maximum
    - address 123 in AS1 generally ≠ address 123 in AS2
    - indirect addressing ~ different ASes to different programs
- ~ *protection between processes*

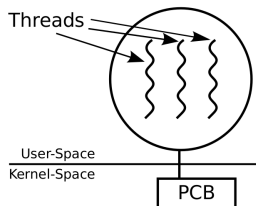
## OS Concepts — Address Space Layout

- **Sections**: address spaces typically laid-out in different sections
  - memory addresses between sections *illegal*
  - illegal addresses ~ page fault (*segmentation fault*)
  - OS usually kills process causing segmentation fault
- **Important sections**:
  - *Stack*: function history, local variables
  - *Data*: Constants, static/global variables, strings
  - *Text*: Program code



## OS Concepts — Threads

- **Thread:** represents execution state of process ( $\geq 1$  thread per process)
  - *IP:* stores currently executed instruction (address in **text** section)
  - *SP:* stores address of stack top ( $> 1$  threads  $\rightarrow$  multiple stacks!)
  - *PSW:* contains flags about execution history (e.g. last calculation was 0  $\rightarrow$  used in following jump instruction)
  - more general purpose registers, floating point registers,...



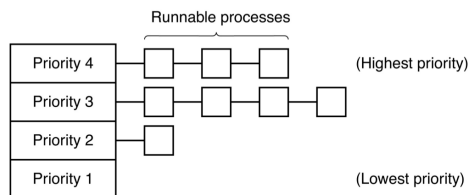
## OS Concepts — Policies vs. Mechanisms

- **Mechanism:** implementation of what is done (e.g. commands to write to HDD)
- **Policy:** rules which decide when what is done and how much (e.g. how often, how many resources are used,...)

$\rightarrow$  mechanisms can be reused even when policy changes

## OS Concepts — Scheduling

- **Motivation:** multiple processes/threads available  $\rightarrow$  OS needs to switch between them (for multitasking)
- **Scheduler:** decides which job to run next (*policy*) — tries to
  - provide fairness
  - meet performance goals
  - adhere to priorities
- **Dispatcher:** performs task-switching (*mechanism*)



## OS Concepts — Files

- **Motivation:** OS hides peculiarities of file storage, programmer uses device-independent *files/directories*
- **Files:** associate *file name* and *offset* with bytes
- **Directories:** associate *directory names* with directory names or file names
- **File System:** ordered block collection
  - main task: translate (dir name + file name + offset) to block
  - programmer uses file system operations to operate on files (**open, read, seek**)
  - processes can communicate directly through special *named pipe* file (used with same operations as any other file)

## OS Concepts — Directory Tree

- **Directories:** form *directory tree/file hierarchy*  $\rightarrow$  structure data
- **Root Directory:** topmost directory in tree
- **Path Name:** used to specify file

## OS Concepts — Mounting

- **Unix:** common to orchestrate multiple file systems in single file hierarchy
- file systems can be *mounted* on directory
- **Win:** manage multiple directory hierarchies with drive letters (e.g. **C:** \Users)

## OS Concepts — Storage Management

- **OS:** provides uniform view of information storage to file systems
  - *Drivers:* hide specific hardware devices  $\rightarrow$  hides device peculiarities
  - general interface abstracts physical properties to logical units  $\rightarrow$  block
- **Performance:** OS increases I/O performance:
  - *Buffering:* Store data temporarily while transferred
  - *Caching:* Store data parts in faster storage
  - *Spooling:* Overlap one job's output with other job's input

### Summary

- **OS:** provides abstractions for and protection between applications
- **Kernel:** does not always run — certain events invoke kernel
  - *syscall:* process asks kernel for service
  - *interrupt:* device sends signal that OS has to handle
  - *exception:* CPU encounters unusual situation
- **Processes:** encapsulate resources needed to run program in OS
  - *threads:* represent different execution states of process
  - *address space:* all memory process can name
  - *resources:* allocated resources, e.g., open files
- **Scheduler** decides which process to run next when multi-tasking
- **Virtual Memory** implements address spaces, provides protection between processes
- **File system** abstracts background store using I/O drivers, provides simple interface (files + directories)

# Processes

## The Process Abstraction

- **Motivation:** computers (seem to) do "several things at the same time" (quick process switching  $\rightarrow$  *multiprogramming*)
- **Model:** *process abstraction* models this concurrency:
  - container contains information about program execution
  - conceptually, every process has own "virtual CPU"
  - execution context is changed on process switch
  - dispatcher switches context when switching processes
  - **context switch:** dispatcher saves current registers/memory mappings, restores those of next process

## Process-Cooking Analogy

- Program/Process like Recipe/Cooking
- **Recipe:** lists ingredients, gives algorithm what to do when
  - $\leadsto$  program describes memory layout/CPU instructions
- **Cooking:** activity of using the recipe
  - $\leadsto$  process is activity of executing a program
- multiple similar recipes for same dish
  - $\leadsto$  multiple programs may solve same problem
- recipe can be cooked in different kitchens at the same time
  - $\leadsto$  program can be run on different CPUs at the same time (as different processes)
- multiple people can cook one recipe
  - $\leadsto$  one process can have several worker threads

## Concurrency vs. Parallelism

- OS uses concurrency + parallelism to implement multiprogramming
  1. **Concurrency:** multiple processes, one CPU
    - $\leadsto$  not at the same time
  2. **Parallelism:** multiple processes, multiple CPU
    - $\leadsto$  at the same time

## Virtual Memory Abstraction — Address Spaces

- every process has own *virtual addresses* (**vaddr**)
- MMU relocates each load/store to *physical memory* (**pmem**)
- processes never see physical memory, can't access it directly
- + MMU can enforce protection (mappings in kernel mode)
- + programs can see more memory than available
  - 80:20 rule: 80% of process memory idle, 20% active
  - can keep working set in RAM, rest on disk
- need special MMU hardware

## Address Space (Process View)

- **Motivation:** code/data/state need to be organized within process
  - $\leadsto$  *address space layout*
- **Data types:**
  1. *fixed size* data items
  2. data naturally *freed in reverse allocation order*
  3. data *allocated/freed "randomly"*
- compiler/architecture determine how large int is and what instructions are used in text section (**code**)

- **Loader** determines based on exe file how executed program is placed in memory

## Segments — Fixed-Size Data + Code

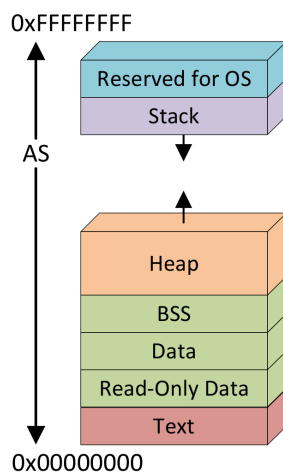
- some data in programs never changes or will be written but never grows/shrinks  
~> memory can be statically allocated on process creation
- **BSS segment** (*block started by symbol*):
  - statically allocated variables/non-initialized variables
  - executable file typically contains starting address + size of BSS
  - entire segment initially 0
- **Data segment**: fixed-size, initialized data elements (e.g. global variables)
- **Read-only data segment**: constant numbers, strings
- All three sometimes summarized as one segment
- compiler and OS decide ultimately where to place which data/how many segments exist

## Segments — Stack

- some data naturally freed in reverse allocation order
  - very easy memory management (stack grows upwards)
- fixed segment starting point
- store top of latest allocation in **stack pointer** (SP) (initialized to starting point)
- *allocate* **a** byte data structure: `SP += a; return(SP - a)`
- *free* **a** byte data structure: `SP -= a`

## Segments — Heap (Dynamic Memory Allocation)

- some data "randomly" allocated/freed
- two-tier memory allocation:
  1. allocate large memory chunk (**heap segment**) from OS
    - base address + **break pointer** (BRK)
    - process can get more/give back memory from/to OS
  2. dynamically partition chunk into smaller allocations
    - `malloc/free` can be used in random order
    - purely user-space, no need to contact kernel



### Summary

**Processes:** recipe vs. cooking = program vs. process

- processes = resource container for OS
- process feels alone (has own CPU and memory)
- OS implements multiprogramming through rapid process switching

# Process API

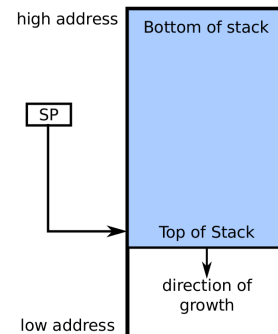
## Execution Model — Assembler (simplified)

- **Principle:** OS interacts directly with compiled programs
  - switch between processes/threads ~> *save/restore* state
  - deal with/pass on *signals/exceptions*
  - receive *requests* from applications
- **Instructions:**
  - `mov`: Copy referenced data from second operand to first operand
  - `add/sub/mul/div`: Add,...from second operand to first operand
  - `inc/dec`: increment/decrement register/memory location

- `shl/shr`: shift first operand left/right by amount given by second operand
- `and/or/xor`: calculate bitwise and,...of two operands storing result in first
- `not`: bitwise negate operand

## Execution Model — Stack (x86)

- **stack pointer** (SP): holds address of stack top (growing downwards)
- **stack frames**: larger stack chunks
- **base pointer** (BP): used to organize stack frames



## Execution Model — jump/branch/call commands (x86)

- `jmp`: continue execution at operand address
- `j$condition`: jump depending on PSW content
  - *true* ~> jump
  - *false* ~> continue
  - examples: `je` (jump equal), `jz` (jump zero)
- `call`: push function to stack and jump to it
- `return`: return from function (jump to return address)

## Execution Model — Application Binary Interface (ABI)

- **Idea:** standardizes binary interface between programs, modules, OS:
  - executable/object file layout
  - calling conventions
  - alignment rules
- **calling conventions:** standardize exact way function calls are implemented  
~> interoperability between compilers

## Execution Model — calling conventions (x86)

- function call — **caller**:
  1. save local scope state
  2. set up parameters where function can find them
  3. transfer control flow
- function call — **called function**:
  1. set up new local scope (local variables)
  2. perform duty
  3. put return value where caller can find it
  4. jump back to caller (IP)

## Passing parameters to the system

- parameters are passed through **system calls**
- call number + specific parameters must be passed
- parameters can be transferred through
  - **CPU registers** (~6)
  - **Main Memory** (heap/stack – more parameters, data types)
- ABI specifies how to pass parameters
- **return code** needs to be returned to application
  - *negative*: error code
  - *positive + 0*: success
  - usually returned via A+D registers

## System call handler

- implements the actual service called through a syscall:
  1. saves tainted registers
  2. reads passed parameters
  3. sanitizes/checks parameters
  4. checks if caller has enough permissions to perform the requested action
  5. performs requested action in behalf of the caller
  6. returns to caller with success/error code

## Process API — creation

- process creation events:
  - system initialization
  - process creation syscall
  - user requests process creation
  - batch job-initiation
- events map to two mechanisms:
  - Kernel spawns initial user space process on boot (Linux: `init`)
  - User space processes can spawn other processes (within their quota)

## Process API — creation (POSIX)

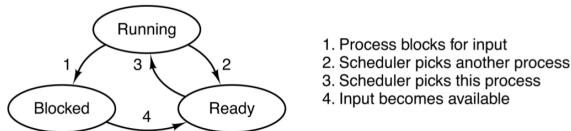
- PID**: identifies process
- pid = fork()**: duplicates current process:
  - returns 0 to new child
  - returns new **PID** to parent→ child and parent independent after `fork`
- exec(name)**: replaces own memory based on executable file
  - name** specifies binary executable file
- exit(status)**: terminates process, returns **status**
- pid = waitpid(pid, &status)**: wait for child termination
  - pid**: process to wait for
  - status**: points to data structure that returns information about the process (e.g., exit status)
  - passed **pid** is returned on success, -1 otherwise
- process tree**: processes create child processes, which create child processes, ...
  - parent and child execute concurrently
  - parent waits for child to terminate (collecting the exit state)

## Daemons

- = program designed to run in the background
- detached from parent process after creation, reattached to process tree root (`init`)

## Process States

- blocking**: process does nothing but wait
  - usually happens on syscalls (OS doesn't run process until event happens)



## Process Termination

- different termination events:
  - normal exit (voluntary)
    - `return 0` at end of `main`
    - `exit(0)`
  - error exit (voluntary)
    - `return x` ( $x \neq 0$ ) at end of `main`
    - `exit(x)` ( $x \neq 0$ )
    - `abort()`
  - fatal error (involuntary)
    - OS kills process after exception
    - process exceeds allowed resources
  - killed by another process (involuntary)
    - another process sends kill signal (only as parent process or administrator)

## Exit Status

- voluntary exit: process returns exit status (integer)
- resources not completely freed after process terminates → **Zombie** or **process stub** (contains exit status until collected via `waitpid`)
- Orphans**: Processes without parents
  - usually adopted by `init`
  - some systems kill all children when parent is killed
- exit status on involuntary exit:
  - Bits 0-6: signal number that killed process (0 on normal exit)
  - Bit 7: set if process was killed by signal
  - Bits 8-15: 0 if killed by signal (exit status on normal exit)

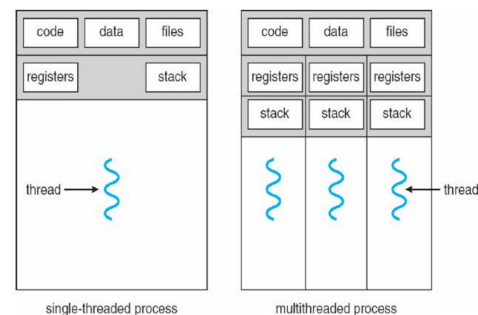
# Threads

## Processes vs. Threads

- Traditional OS**: each process has
  - own address space
  - own set of allocated resources
  - one thread of execution (= one execution state)
- Modern OS**: processes + threads of execution handled more flexibly
  - processes* provide abstraction of address space and resources
  - threads* provide abstraction of execution states of that address space
- Exceptions**:
  - sometimes different threads have different address spaces
  - Linux: threads = regular processes with shared resources and AS regions

## Threads — why?

- many programs do multiple things at once (e.g. web server)
  - writing program as many sequential threads may be easier than with blocking operations
- Processes**: rarely share data (if, then explicitly)
- Threads**: closely related, share data



## Threads — POSIX

- PThread**: base object with
  - identifier* (thread ID, TID)
  - register set* (including IP and SP)
  - stack area* to hold execution state
- Pthread\_create**: create new thread
  - Pass: *pointer* to `pthread_t` (will hold TID after successful call)
  - Pass: *attributes, start function, arguments*
  - Returns: 0 on success, error value else
- Pthread\_exit**: terminate calling thread
  - Pass: exit code (casted to void pointer)
  - Free's resources (e.g. stack)
- Pthread\_join**: wait for specified thread to exit
  - Pass: `pthread_t` to wait for (or -1 for any thread)
  - Pass: pointer to pointer for exit code
  - Returns: 0 on success, error value else
- Pthread\_yield**: release CPU to let another thread run

## Threads — Problems

- Processes vs. Threads**:
  - Processes*: only share resources explicitly
  - Threads*: more shared state → more can go wrong
- Challenges**: programmer needs to take care of
  - activities*: dividing, ordering, balancing
  - data*: dividing
  - shared data*: access synchronizing

## PCB vs. TCP

- PCB (process control block)**: information needed to implement processes
  - always known to OS
- TCB (thread control block)**: per thread data
  - OS knowledge depends on *thread model*



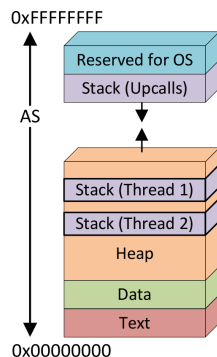
PCB	TCB
Address space	Instruction pointer
Global variables	Registers
Open files	Stack
Child processes	State
Pending alarms	

## Thread models

- **Kernel Thread:** known to OS kernel
- **User Thread:** known to process
- **N:1-Model:** kernel only knows one of possibly multiple threads
  - N:1 user threads = *user level threads* (ULT)
- **1:1-Model:** each user thread maps to one kernel thread
  - 1:1 user threads = *kernel level threads* (KLT)
- **M:N-Model** (hybrid model): flexible mapping of user threads to less kernel threads

## Thread models — N:1

- Kernel only manages process → multiple threads unknown to kernel
- Threads managed in user-space library (e.g. GNU Portable Threads)
- **Pro:**
  - + faster thread management operations (up to 100 times)
  - + flexible scheduling policy
  - + few system resources
  - + usable even if OS doesn't support threads
- **Con:**
  - no parallel execution
  - whole process blocks if one user thread blocks
  - reimplementing OS parts (e.g. scheduler)
- **Stack:**
  - main stack known to OS used by thread library
  - own execution state (= stack) dynamically allocated by user thread library for each thread
  - possibly own stack for each exception handler
- **Heap:**
  - concurrent heap use possible
  - *Attention:* not all heaps are reentrant
- **Data:** divided into BSS, data and read-only data here as well

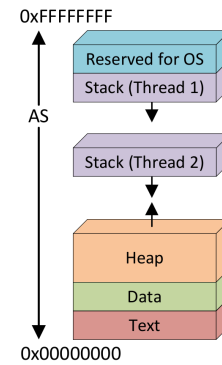


## Thread models — 1:1

- kernel knows + manages every thread
- **Pros:**
  - + real parallelism possible
  - + threads block individually
- **Cons:**
  - OS manages every thread in system (TCB, stacks,...)
  - Syscalls needed for thread management
  - scheduling fixed in OS
- **Stack:**
  - own execution state (= stack) for every thread
  - possibly own stack for (each) exception handler
- **Heap:**
  - parallel heap use possible
  - *Attention:* not all heaps are thread-safe
  - if thread-safe: not all heap implementations perform well with many threads
- **Data:** divided into BSS, data and read-only data here as well

## Thread models — M:N

- **Principle:**  $M$  ULTs are maps to (at most)  $N$  KLT
  - *Goal:* pros of ULT and KLT — non-blocking with quick management
  - create sufficient number of KLTs and flexibly allocate ULTs to them



- *Idea:* if ULT blocks ULTs can be switched in userspace
- **Pros:**
  - + flexible scheduling policy
  - + efficient execution
- **Cons:**
  - hard to debug
  - hard to implement (e.g. blocking, number of KLTs,...)
- **Implementation — Up-calls:**
  - kernel notices that thread will block → sends signal to process
  - up-call notifies process of thread id and event that happened
  - exception handler of process schedules a different process thread
  - kernel later informs process that blocking event finished via other up-call

## Summary

- programs often do closely related things at once
    - mapped to thread abstraction: multiple threads of execution operate in same process
  - differentiation between process information (PCB) and thread information (TCB)
  - **thread models:**
    - $N : 1$ : threads fully managed in user-space
    - $1 : 1$ : threads fully managed by kernel
    - $M : N$ : threads are flexibly managed either in user-space or kernel
  - multi-threaded programs operate on same data concurrently or even parallel:
    - *synchronization:* accessing such data must be synchronized
- makes writing such programs challenging

# Scheduling

## Motivation

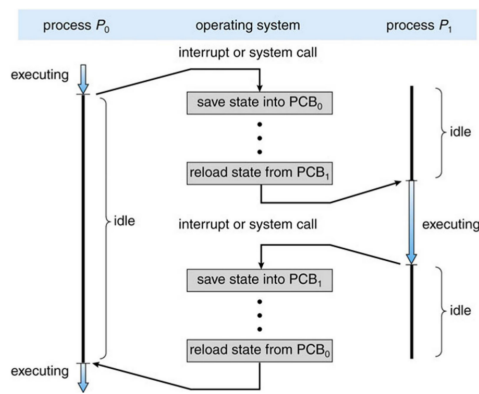
- $K$  jobs ready to run,  $K > N \geq 1$  CPUs available
- **Scheduling Problem:**
  - Which jobs should kernel assign to which CPUs?
  - When should it make decision?

## Dispatcher

- **Dispatcher:** performs actual process switch
  - mechanism
  - save/restore process context
  - switch to user mode
- **Scheduler:** selects next process to run based on *policy*

## Voluntary Yielding vs. Preemption

- kernel responsible for CPU switch
- kernel doesn't always run → can only dispatch different process when invoked
- **cooperative multitasking:** running process performs *yield* syscall
  - kernel switches process
- **preemptive scheduling:**
  - kernel invoked in certain time intervals
  - kernel makes scheduling decisions after every time-slice



## Scheduling — Process States

- **new:** process was created but did not run yet
- **running:** instructions are currently being executed
- **waiting:** process is waiting for some event
- **ready:** process is waiting to be assigned a processor
- **terminated:** process has finished execution

## Scheduling — long-term vs. short-term

- **Short-term scheduler** (CPU Scheduler, focused on in this lecture):
  - selects process to run next, allocates CPU
  - invoked frequently (ms)  $\leadsto$  must be fast
- **Long-term scheduler** (job scheduler):
  - selects process to be brought into ready queue
  - invoked very infrequently (s, m)  $\leadsto$  can be slow
  - controls degree of *multiprogramming*

## Scheduling queues

- **job queue:** set of all processes in system
- **ready queue:** process in main memory, ready or waiting
- **device queue:** processes waiting for I/O device

## Scheduling Policies — Categories

- **batch scheduling:**
  - still widespread in business (payroll, inventory,...)
  - no users waiting for quick response
  - non-preemptive algorithms acceptable  $\rightarrow$  less switches  $\rightarrow$  less overhead
- **interactive scheduling:**
  - need to optimize for response time
  - preemption essential to keep processes from hogging CPU
- **real-time scheduling:**
  - guarantee job completion within time constraints
  - need to be able to plan when which process runs + how long
  - preemption not always needed

## Scheduling Policies — Goals

- **General:**
  - *fairness:* give each process fair share of CPU
  - *balance:* keep all parts of system busy
- **batch scheduling:**
  - *throughput:* number of processes that complete per time unit
  - *turnaround time:* time from job submission to job completion
  - *CPU utilization:* keep CPU as busy as possible
- **interactive scheduling:**
  - *waiting time:* reduce time a process waits in waiting queue
  - *response time:* time from request to first response
- **real-time scheduling:**
  - *meeting deadlines:* finishing jobs in time
  - *predictability:* minimize jitter

## Scheduling Policies — first come first served

- intuitively clear
- **Example:** 3 processes arrive at time 0 in the order  $P_1, P_2, P_3$

Process	Burst time	Turnaround time
$P_1$	24	24
$P_2$	3	27
$P_3$	3	30

- $\leadsto$  average turnaround time 27  $\rightarrow$  can we do better?
- **Conclusion:** if processes would arrive in order  $P_2, P_3, P_1$ , average turnaround time would be 13
- $\leadsto$  good scheduling can reduce turnaround time

## Scheduling Policies — shortest job first

- **Benefits:** optimal average turnaround/waiting/response time
- **Challenge:** cannot know job lengths in advance
- **Solution:** predict length of next CPU burst for each process
  - $\leadsto$  schedule process with shortest burst next
- **Burst Estimation:** *exponential averaging*
  - $\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n$
  - ( $t_n$ : actual length of  $n$ -th CPU burst,  $\tau_{n+1}$ : predicted length of next CPU burst,  $0 \leq \alpha \leq 1$ )

## Process Behavior — CPU bursts

- CPU bursts exist because processes wait for I/O
- **CPU-bound processes:** spends more time doing computations
  - $\leadsto$  few very long CPU bursts
- **I/O-bound processes:** spends more time doing I/O
  - $\leadsto$  many short CPU bursts

## Scheduling Policies — preemptive shortest-job-first

- SJF optimizes waiting/response time
  - $\leadsto$  what about throughput?
- **Problem:** CPU-bound jobs hold CPU until exit or I/O  $\rightarrow$  poor I/O utilization
- **Idea:** SJF, but preempt periodically to make new scheduling decision
  - each time slice: schedule job with shortest remaining time next
  - alternatively: schedule job with shortest next CPU burst

## Scheduling Policies — round robin

- **Problem:** batch schedulers suffer from starvation and don't provide fairness
- **Idea:** each process runs for small CPU time unit
  - *time quantum/time slice* length: usually 10-100ms
  - preempt processes that have not blocked by end of time slice
  - append current thread to end of run queue, run next thread
- **Caution:** time slice length needs to balance interactivity and overhead!
  - $\rightarrow$  if time slice length in the area of dispatch time, 50% of CPU time wasted for process switching

## Scheduling Policies — virtual round robin

- **Problem:** RR is unfair for I/O-bound jobs: they block before using up time quantum
- **Idea:** put jobs that didn't use up their quantum in additional queue
  - store share of unused time-slice
  - give those jobs additional queue priority
  - put them back into normal queue afterwards

## Scheduling Policies — (strict) priority scheduling

- **Problem:** not all jobs are equally important
  - $\leadsto$  different priorities (e.g., 4)
- **Solution:** associate priority number with each process
  - RR for each priority
  - *aging:* old low priority processes get executed before new higher priority processes

## Scheduling Policies — multi-level feedback queue

- **Problem:** context switching expensive
  - $\leadsto$  trade-off between interactivity and overhead?
- **Goals:**
  - higher priority for I/O jobs (usually don't use up quantum)
  - low priority for CPU jobs (rather run them longer)
- **Idea:** different queues with different priorities and time slice lengths
  - schedule queues with (static) priority scheduling
  - double time slice length in each next-lower priority
  - process to higher priority when they don't use up quantum repetitively
  - process to lower priority when they use up quantum repetitively

## Scheduling Principles — priority donation

- **Problem:** Process B (higher priority) waits for process A (lower priority)  
→ B has now effectively lower priority
- **Solution:** *priority donation*
  - give A priority of B as long as B waits for A
  - if C, D, E wait for B → A gets highest priority of B, C, D, E

## Scheduling Policies — lottery scheduling

- issue number of lottery tickets to processes (amount depending on priority)
- amount of tickets controls average proportion of CPU for each process
- **Scheduling:** scheduler draws random number  $N$ , process with  $N$ -th ticket is executed
- processes can transfer tickets to other processes if they wait for them

### Summary

- **phases:** processes have phases of communication and waiting for I/O  
→ appropriate switching between processes increases computing system utilization
- **goal-based:** scheduler decides what appropriate means based on goals
  - *long-term scheduler:* degree of multiprogramming
  - *short-term scheduler:* which process to run next
- **dispatching:** only happens when OS is invoked
  - *cooperative scheduling:* currently running thread yields (syscall)
  - *preemptive scheduling:* OS is called periodically to switch threads

# Inter Process Communication

## Overview

- **Reasons** for cooperating processes:
  - *information sharing:* share file/data-structure in memory
  - *computation speed-up:* break large tasks in subtasks → parallel execution
  - *modularity:* divide system into collaborating modules with clean interfaces
- **IPC:** allows data exchange
  - *message passing:* explicitly send/receive information using syscalls
  - *shared memory:* physical memory region used by multiple processes/threads

## IPC — message passing

- = mechanism for processes to communicate and synchronize
- message passing facilities generally provide **send** and **receive**
- **Implementations:**
  - hardware bus
  - shared memory
  - kernel memory
  - network interface card (NIC)
- **Direct messages:** processes explicitly named when exchanging messages
- **Indirect messages:** sending to/receiving from *mailboxes*
  - first communicating process creates mailbox, last destroys
  - processes can only communicate through shared mailbox

## Indirect messages – synchronization

- **Blocking** (synchronous):
  - *blocking send:* sender blocks until message is received
  - *blocking receive:* receiver blocks until message is available
- **Non-blocking** (asynchronous):
  - *non-blocking send:* sender sends message, then continues
  - *non-blocking receive:* receiver receives valid message or **null**

## Messaging — Buffering

- messages are *queued* using different capacities while being in-flight
- **zero capacity:** no queuing
  - *rendezvous:* sender must wait for receiver
  - message is transferred as soon as receiver becomes available → no latency/jitter
- **bounded capacity:** finite number + length of messages
  - sender can send before receiver waits for messages
  - sender must wait if link is full
- **unbounded capacity:**
  - sender never waits

- memory may overflow → potentially large latency/jitter between **send** and **receive**

## Messaging — POSIX message queues

- **create** or open existing message queue:  
`mqd_t mq_open (const char *name, int oflag);`
  - **name** is path in file system
  - access permission controlled through file system access permission
- **send** message to message queue:  
`int mq_send (mqd_t md, const char *msg, size_t len, unsigned priority);`
- **receive** message with highest priority in message queue:  
`int mq_receive (mqd_t md, char *msg, size_t len, unsigned *priority);`
- **register** callback handler on message queue (to avoid polling):  
`int mq_notify (mqd_t md, const struct sigevent *sevp);`
- **remove** message queue:  
`int mq_unlink (const char *name);`

## Shared Memory

- **Principle:** communicate through region of shared memory
  - every write to shared region is visible to all other processes
  - hardware guarantees that always most recent write is read
- **Implementation:** message passing via shared memory is application-specific
- **Problems:** using shared memory in a safe way is tricky
  - *cache coherency protocol:* makes usage with many processes/CPU hard
  - *race conditions:* makes usage with multiple writers hard

## Shared Memory — POSIX shared memory

- **create** or open existing POSIX shared memory object:  
`int shm_open (const char *name, int oflag, mode_t mode);`
- **set** size of shared memory region:  
`ftruncate (smd, size_t len);`
- **map** shared memory object to address space:  
`void* mmap (void* addr, size_t len, [...], smd, [...]);`
- **unmap** shared memory object from address space:  
`int munmap (void* addr, size_t len);`
- **destroy** shared memory object:  
`int shm_unlink (const char *name);`

## Shared Memory — sequential memory consistency

- = *the result of execution as if all operations were executed in some sequential order, and the operations of each processor occurred in the order specified by the program.*
- **Model:**
  - all memory operations occur one at a time in *program order*
  - ensures write atomicity
- **Reality:** compiler and CPU re-order instructions to *execution order*  
→ without SC many processes on many CPU behave worse than preemptive threads on 1 CPU

## Shared Memory — memory consistency model

- **Problem:**
  - CPUs generally not sequentially consistent
  - compilers do not generate code in program order

## Synchronization — race conditions

- **Assume:** sequential memory consistency → no atomic memory transactions!
- **Critical Sections:** protect instructions inside critical section from concurrent execution

## Critical Sections — desired properties

- **mutual exclusion:** at most one thread can be in the CS at any time
- **progress:** no thread running outside of CS may block other thread from getting in
- **bounded waiting:** once a thread starts trying to enter CS, there is a bound on number of times other threads get in

## Critical Sections — disabling interrupts

- kernel only switches on interrupts (usually on *timer interrupt*)  
→ have per-thread *do not interrupt* (DNI)-bit
- **single-core system:**



- enter CS: set DNI bit
- leave CS: clear DNI bit
- **Advantages:**
  - + easy + convenient in kernel
- **Disadvantages:**
  - *only works on single-core systems*: disabling interrupts on one CPU doesn't affect other CPUs
  - *only feasible in kernel*: don't want to give user power to turn off interrupts!

## Critical Sections — lock variables

- define global **lock** variable
  - only enter CS if **lock** is 0, set to 1 on enter
  - wait for lock to become 0 otherwise (*busy waiting*)
- **Problem:** doesn't solve CS problem! Reading/Setting lock not atomic!

## Critical Sections — spinlocks

- to make lock variable approach work, lock variable must be tested and set at same time atomically:
- **x86: xchg** can atomically exchange memory content with register
  - exchanges register content with memory content
  - returns previous memory content of lock
- implementation of critical section as *spinlock*:

```
void enter_critical_section (volatile bool *lock) {
    while (xchg(lock, 1) == 1); // lock = 1, return old value
                                // repeat until old value != 1
}

void leave_critical_section (volatile bool *lock) {
    *lock = 0;
}
```

- **Advantages:**
  - + *mutual exclusion*: only one thread can enter CS
  - + *progress*: only thread within CS hinders others of getting in
- **Disadvantages:**
  - *bounded waiting*: no upper bound

## Spinlocks — Limitations

- **Congestion:**
  - if most times there is no thread in CS when another tries to enter, then spinlocks are very easy + efficient
  - if CS is large or many threads try to enter, spinlocks might not be good choice as all threads actively wait spinning
- **Multicore:** memory address is written at every atomic swap operation
  - memory is expensively kept coherent
- **Static Priorities** (e.g., *priority inversion*): if low-priority threads hold lock it will never be able to release it, because it will never be scheduled

## Spinlocks — sleep while wait

- **Problem:** busy part of busy waiting
  - wastes resources,
  - stresses cache coherence protocol,
  - can cause priority inversion problem
- **Idea:**
  - threads sleep on locks if occupied
  - wake up threads one at a time when lock becomes free

## Spinlocks — semaphore

- two new syscalls operating on **int** variables:
  - **wait (&s)**: if **s** > 0: **s--** and continue, otherwise let caller sleep
  - **signal (&s)**: if no thread is waiting: **s++**, otherwise wake one up
- initialize **s** to maximum number of threads that may enter CS
  - **wait** = **enter\_critical\_section()**
  - **signal** = **leave\_critical\_section()**
- **mutex** (semaphore): semaphore initialized to 1 (only admits one thread at a time into CS)
- **counting semaphore**: semaphore allowing more than one thread into CS at a time

## Semaphore — implementation

- **wait** and **signal** calls need to be carefully synchronized (otherwise *race condition* between checking and decrementing **s**)
- **signal loss** can occur when waiting and waking threads up at same time

- each semaphore has **wake-up queue**:
  - *weak semaphores*: wake up random waiting thread on **signal**
  - *strong semaphores*: wake up thread strictly in order which they started **waiting**
- **Advantages:**
  - + *mutual exclusion*: only one thread can enter CS for mutexes
  - + *progress*: only thread within CS hinders others to get in
  - + *bounded waiting*: strong semaphores guarantee bounded waiting
- **Disadvantages:**
  - every enter and exit of CS is syscall → slow

## Fast User Space mutex

- **spinlock:**
  - + quick when wait-time is short
  - waste resources when wait-time is long
- **semaphore:**
  - + efficient when wait-time is long
  - syscall overhead at every operation
- **futex:**
  - userspace + kernel component
  - try to get into CS with userspace spinlock
    - CS busy → use syscall to put thread to sleep
    - otherwise → enter CS with now locked spinlock completely in userspace

### Summary

- **communication** between processes/threads often needed
  - *message passing*: provide explicit send/receive functions to exchange messages
  - *implicitly/explicitly shared memory* between threads/processes: allows information exchange
- **data races**: need to be taken into account when communicating
- **synchronization techniques:**
  - interlocked atomic operations
  - spinlocks
  - semaphores
  - futexes

# Synchronization and Deadlocks

## Producer-Consumer Problem

- **Definition:**
  - buffer is shared between producer and consumer (LIFO)
  - **count** integer keeps track of number of currently available items
  - producer produces item → placed in buffer, **count++**
  - buffer full → producer needs to sleep until consumer consumed an item
  - consumer consumes item → remove item from buffer, **count--**
  - buffer empty → consumer needs to sleep until producer produces item
- **Problem:** *race condition* on **count**

## Producer-Consumer Problem — condition variables

- **Solution:** can be solved with mutex + 2 counting semaphores
  - hard to understand
  - hard to get right
  - hard to transfer to other problems
- **condition variables:** allow blocking until condition is met
  - usually suitable for same problems but much easier to get right
- **Idea:**
  - new operation performs *unlock, sleep, lock* atomically
  - new wake-up operation is called with lock held
- simple mutex lock/unlock around CS + no signal loss
- **Pthread** condition variables:
  - **pthread\_cond\_init**: create + initialize new CV
  - **pthread\_cond\_destroy**: destroy + free existing CV
  - **pthread\_cond\_wait**: block waiting for signal
  - **pthread\_cond\_timedwait**: block waiting for signal or timer
  - **pthread\_cond\_signal**: signal another thread to wake up
  - **pthread\_cond\_broadcast**: signal all threads to wake up

## Reader-Writer Problem

- **Problem:** model access to shared data structures

```

void producer()
{
    Item newItem;
    for(;;) // ever
    {
        newItem = produce();
        mutex_lock( &lock );
        while( count == MAX_ITEMS )
            cond_wait( &less, &lock );

        insert( newItem );
        count++;
        cond_signal( &more );
        mutex_unlock( &lock );
    }
}

void consumer()
{
    Item item;
    for(;;) // ever
    {
        mutex_lock( &lock );
        while( count == 0 )
            cond_wait( &more, &lock );

        item = remove();
        count--;
        cond_signal( &less );
        mutex_unlock( &lock );
        consume( item );
    }
}

```

- many threads compete to read/write same data
- readers: only read data set, not performing any updates
- writers: both read and write
- using single mutex for read/write operations is not a good solution! (unnecessarily blocking out multiple readers while no writer is present)
- Idea:** locking should reflect different semantics for reading/writing
  - no writing thread → multiple readers may be present
  - writing thread → no other reader/writer allowed

## Dining-Philosophers Problem

- Definition:** 5 philosophers with cyclic workflow:
  - think
  - get hungry
  - grab one chopstick
  - grab other chopstick
  - put down chopsticks
- Rules:**
  - no communication
  - no atomic grabbing of both chopsticks
  - no wrestling
- Abstraction:** models threads competing for limited number of resources **Problem:** what happens if all philosophers grab left chopstick at once?
- Deadlock workarounds:**
  - deadlock avoidance: just 4 philosophers allowed at table of 5
  - deadlock prevention: odd philosophers take left chopstick first, even ones take right first → deadlock prevention



## Deadlocks

- Deadlocks** can arise if all four conditions hold simultaneously:
  - mutual exclusion:** limited resource access (can only be shared with finite number of users)
  - hold and wait:** wait for next resource while already holding at least one
  - no preemption:** granted resource cannot be taken away but only handed back voluntarily
  - circular wait:** possibility of circularity in requests graph

## Deadlocks — countermeasures

- prevention:** pro-active, make deadlocks impossible to occur
- avoidance:** decide on allowed actions based on a-priori knowledge
- detection (recovery):** react after deadlock happened

## Deadlocks — prevention

- Goal:** negate at least one of the required deadlock conditions:
  - mutual exclusion:** buy more resources, split into pieces, virtualize
  - hold and wait:** get all resources en-bloque, 2-phase-locking
  - no preemption:** virtualize to make preemptable
  - circular wait:** reorder resources, prevent through partial order on resources


## Deadlocks — avoidance

- safe state:** system is in safe state → no deadlocks

- unsafe state:** system is in unsafe state → deadlocks possible
- avoidance:** on every resource request decide if system stays in safe state → *resource allocation graph*


## Deadlock Avoidance — resource allocation graph

- principle:** view system state as graph
  - processes* = round nodes
  - resources* = square nodes
  - resource instance* = dot in resource node
  - resource requests/assignments* = edges
    - resource → process = resource is assigned to process
    - process → resource = process is requesting resource



## Deadlocks — detection

- Principle:** allow system to enter deadlock → detection → recovery scheme
- wait-for graph (WFG):**
  - processes* = nodes
  - wait-for relationship* = edges
- periodically invoke algorithm searching for cycle in graph
  - cycle exists → deadlock exists



## Deadlocks — recovery

- Process termination:**
  - all:* abort all deadlocked processes
  - selective:* abort one process at a time until deadlock is eliminated
- Termination order:** in which order should processes be aborted?
  - process priority
  - how long already computed? how much longer for completion?
  - amount of resources used
  - amount of resources needed for completion
  - how many processes will need to be terminated
  - interactive or batch?
- Resource preemption:**
  - victim selection:* minimize cost
  - rollback:* perform periodic snapshots, abort process to preempt resources → restart from last safe state
  - starvation:* same process may always be picked as victim → include rollback count in cost factor

### Summary

- classical synchronization problems:** model synchronization problems occurring in reality
  - producer-consumer:* shared use of buffers/queues
  - reader-writer:* shared access to data structures
  - dining philosophers:* competition for limited resources
- such synchronization problems occur very often when programming operating systems
- parallelism:** introduced by multiple processors + multiprogramming, needs to be considered carefully when writing OS