

## I. INTRO TO C

### Hello World

```
#include <stdio.h>
```

```
int main(void) {
    printf("Hello World!\n");
    return 0;
}
```

**#include:** preprocessor inserts `stdio.h` contents

**stdio.h:** contains `printf` declaration

**main:** program starts here

**void:** keyword for argument absence

**{ }:** basic block/scope delimiters

**printf:** prints to the terminal

**\n:** newline character

**return:** leave function, return value

### Compiling

```
$ gcc hello.c -o hello
$ ./hello
Hello World!
```

### Basic Data Types

```
char c = 5; char c = 'a';
    one byte, usually for characters (1970: ASCII is fine)
int i = 5; int i = 0xf; int i = 'a';
    usually 4 bytes, holds integers
float f = 5; float f = 5.5;
    4 bytes floating point number
double d = 5.19562
    8 bytes double precision floating point number
```

### Basic Data Types – logic

```
int i = 5 / 2; //i = 2
    integer logic, no rounding
float f = 5.0f / 2; //f = 2.5f
    decimal logic for float and double
char a = 'a' / 2 //a = 97 / 2 = 48
    char interpreted as character by console
```

### Basic Data Types – signed/unsigned

```
signed int i = -5 //i = -5 (two's complement)
unsigned int i = -5 //i = 4294967291
```

### Basic Data Types – short/long

```
short int i = 1024 //-32768...32767
long int i = 1024 //-2147483648...2147483647
```

### Basic Data Types – more size stuff

```
sizeof int; sizeof long int; //4; 4; (x86 32-Bit)
    use data types from inttypes.h to be sure about sizes:

#include <inttypes.h>
int8_t i; uint32_t j;
```

### Basic Data Types – const/volatile

```
const int c = 5;
    i is constant, changing it will raise compiler error
volatile int i = 5;
    i is volatile, may be modified elsewhere (by different program
    in shared memory, important for CPU caches, register, assump-
    tions thereof)
```

### Variables – local vs. global

```
int m; // global variable

int myroutine(int j) {
    int i = 5 // local variable
    i = i+j;
    return i;
}
```

**global variables (`int m`):**

lifetime: while program runs

placed on pre-defined place in memory

**basic block/function-local variables (`int i`):**

lifetime: during invocation of routine

placed on stack or in registers

### Variables – local vs. static

```
int myroutine(int j) {
    static int i = 5;
    i = i+j;
    return i;
}
```

```
k = myroutine(1); // k = 6
```

```
k = myroutine(1); // k = 7
```

**static function-local variables:**

saved like global variables

variable persistent across invocations

lifetime: like global variables

### Printing

```
int i = 5; float f = 2.5;
printf("The numbers are i=%d, f=%f", i, f);
```

comprised of format string and arguments

may contain format identifiers (`%d`)

see also [man printf](#)

**special characters:** encoded via leading backslash:

```
\n newline
\t tab
\' single quote
\" double quote
\0 null, end of string
```

### Compound data types

**structure:** collection of named variables (different types)

**union:** *single* variable that can have multiple types

members accessed via `.` operator

```
struct coordinate {
    int x;
    int y;
}
```

```
union longorfloat {
    long l;
    float f;
}
```

```
struct coordinate c;
c.x = 5;
c.y = 6;
```

```
union longorfloat lf;
lf.l = 5;
lf.f = 6.192;
```

## Functions

encapsulate functionality (*reuse*)  
 code structuring (*reduce complexity*)  
 must be **declared** and **defined**  
Declaration: states signature  
Definition: states implementation (implicitly declares function)

```
int sum(int a, int b); // declaration
```

```
int sum(int a, int b) { // definition
    return a+b;
}
```

## Header files

header file for frequently used declarations  
 use **extern** to declare global variables defined elsewhere  
 use **static** to limit scope to current file (e.g. **static float pi** in **sum.c**: no **pi** in **main.c**)

```
// mymath.h
int sum(int a, int b);
extern float pi;
```

```
// sum.c
#include "mymath.h"

float pi = 3.1415927;
int sum(int a, int b) {
    return a+b;
}
```

```
// main.c
#include <stdio.h>
#include "mymath.h"

void main() {
    printf("%d\n", sum(1,2));
    printf("%f\n", pi);
}
```

## Data Segments and Variables

Stack: local variables  
Heap: variables created at runtime via **malloc()**/**free()**  
Data Segment: static/global variables  
Code: functions

## Function overloading

*no function overloading in C!*  
 use arrays or pointers

## Pointers

```
int a = 5;
int *p = &a // points to int, initialized to point to a
int *q = 32 // points to int at address 32
int b = a+1;
int c = *p; // dereference(p) = dereference(&a) = 5
int d = (*p)+2 // = 7
int *r = p+1; // pointing to next element p is pointing to
int e = *(p+2) // dereference (p+2) = d = 7
```

## Pointers – linked list

linked-list implementation via next-pointer

```
struct ll {
    int item;
    struct ll *next;
}

struct ll first;
first.item = 123;

struct ll second;
second.item = 456;
first.next = &second;
```

## Arrays

= fixed number of variables *continuously laid out in memory*

```
int A[5]; // declare array (reserve memory space)
A[4] = 25; A[0] = 24; // assign 25 to last, 24 to first elem
char c[] = {'a',5,6,7,'B'} // init array, length implicit
c[64] = 'Z' // NO bounds checking at compile/run (may raise
              protection fault)
```

// declare pointer to array; address elements via pointer:

```
char *p = c;
*(p+1) = 'Z'; p[3] = 'B'; char b = *p; // = 'a'
```

## Strings

= array of **chars** terminated by **NULL**:

```
char A[] = { 'T', 'e', 's', 't', '\0' };
char A[] = "Test";
```

declaration via pointer:

```
const char *p = "Test";
```

common string functions (**string.h**):

length: **size\_t strlen(const char \*s, size\_t maxlen)**

compare:

```
int strcmp(const char *s1, const char *s2, size_t n);
```

```
copy: int strcpy(char *dest, const char *src, size_t n);
```

```
tokenize: char *strtok(char *str, const char *delim);
```

(e.g. split line into words)

## Arithmetic/bitwise operators

arithmetic operators:

```
a+b, a++, ++a, a+=b, a-b, a--, --a, a-=b, a*b, a*=b, a/b, a/=b, a%b,
a%=b
```

logical operators:

```
a&b, a|b, a>>b, a<<b, a^b, ~a
```

difference pre-/post-increment:

```
int a = 5;
if(a++ == 5) printf("Yes"); // Yes
a = 5;
if(++a == 5) printf("Yes"); // nothing
```

operators in order of precedence:

```
( ), [ ], -, .
!, ++, --, +y, -y, *z, &=, (type), sizeof
*, /, %
+, -
<<, >>
<, <=, >, >=
==, !=
&
^
|
&&
||
?, :
=, +=, -=, *=, /=, %=, &=, ~=, =, <=, >=
,
```

## Structures

brackets only needed for multiple statements

**if/else**, **for**, **while**, **do-while**, **switch**

may use **break/continue**

**switch**: need **break** statement, otherwise will fall through

```
if(a==b) printf("Equal") else printf("Different");
for(i=10; i>=10; i--) printf("%d", i+1);
int i=10; while(i-->0) printf("foo");
int i=0; do printf("bar"); while(i++ != 0);
```

```
char a = read();
switch(a) {
    case '1':
        handle_1();
        break;
    default:
        handle_other();
        break;
}
```

## Type casting

explicit casting: precision loss possible  
`int i = 5; float f = (float)i;`  
 implicit casting: if no precision is lost  
`char c = 5; int i = c;`  
 pointer casting: changes address calculation  
`int i = 5; char *p = (char *)&i; *(p+1) = 5;`  
 type hierarchy: „wider“/„shorter“ types  
`unsigned int` wider than `signed int`  
 operators cast parameters to widest type  
 Attention: assignment cast after operator cast

## C Preprocessor

modifies *source code* before compilation  
 based on preprocessor *directives* (usually starting with `#`)  
`#include <stdio.h>, #include "mystdio.h":`  
 copies contents of file to current file  
 only works with strings in source file  
 completely ignores C semantics

### Preprocessor – search paths

`#include <file>`: system include, searches in:  
`/usr/local/include`  
`libdir/gcc/[target]/[version]/include`  
`/usr/[target]/include`  
`/usr/include`  
 (target: arch-specific (e.g. i686-linux-gnu),  
 version: gcc version (e.g. 4.2.4))  
`#include "file"`: local include, searches in:  
 directory containing current file  
 then paths specified by `-i <dir>`  
 then in system include paths

### Preprocessor – definitions

defines introduce replacement strings (can have arguments, based on string replacement)  
 can help code structuring, often leading to source code cluttering

```
#define PI 3.14159265
#define TRUE (1)
#define max(a,b) ((a > b) ? (a) : (b))
#define panic(str) do { printf(str); for (;;) } while(0);

#ifdef __unix__
#include <unistd.h>
#elif defined _WIN32
#include <windows.h>
#endif
```

### Preprocessor – predefined macros

system-specific:  
`__unix__, __WIN32__, __STDC_VERSION__`  
 useful:  
`__LINE__, __FILE__, __DATE__`

## Libraries

= collection of functions contained in object files, glued together in dynamic/static library  
 ex.: Math header contains declarations, but not all definitions  
 ~ need to link math library: `gcc math.c -o math -lm`

```
#include <math.h>
#include <stdio.h>

int main() {
    float f = 0.555f;
    printf("%f", sqrt(f*4));
    return 0;
}
```

## II. INTRODUCTION TO OPERATING SYSTEMS

### What's an OS?

**abstraction**: provides abstraction for applications  
 manages and hides hardware details  
 uses low-level interfaces (not available to applications)  
 multiplexes hardware to multiple programs (*virtualisation*)  
 makes hardware use efficient for applications

### protection:

from processes using up all resources (*accounting, allocation*)  
 from processes writing into other processes memory

### resource managing:

manages + multiplexes hardware resources  
 decides between conflicting requests for resource use  
 strives for efficient + fair resource use

### control:

controls program execution  
 prevents errors and improper computer use


~ no universally accepted definition

### Hardware Overview

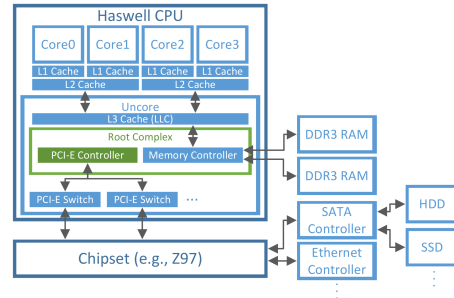
CPU(s)/devices/memory (conceptually) connected to common bus

CPU(s)/devices competing for memory cycles/bus

all entities run concurrently



today: multiple busses



### Central Processing Unit (CPU) – Operation

fetches instructions from memory, executes them (instruction format/-set depends on CPU)

CPU internal registers store (meta-)data during execution (general purpose registers, floating point registers, instruction pointer (IP), stack pointer (SP), program status word (PSW),...)


### execution modes:

**user mode** (x86: *Ring 3/CPL 3*):

only non-privileged instructions may be executed  
 cannot manage hardware → **protection**

**kernel mode** (x86: *Ring 0/CPL 0*):

all instructions allowed  
 can manage hw with **privileged instructions**



## Random Access Memory (RAM)

keeps currently executed instructions + data  
 today: CPUs have built-in *memory controller*  
 root complex connected directly via  
 „wire“ to caches  
 pins to RAM  
 pins to PCI-E switches

## Caching

RAM delivers instructions/data slower than CPU can execute  
 memory references typically follow *locality principle*:  
**spatial locality**: future refs often near previous accesses  
 (e.g. next byte in array)  
**temporal locality**: future refs often at previously accessed ref  
 (e.g. loop counter)  
*caching* helps mitigating this memory wall:  
 copy used information temporarily from slower to faster storage  
 check faster storage first before going down **memory hierarchy**  
 if not, data is copied to cache and used from there

### Access latency:

register:  $\sim 1$  CPU cycle  
 L1 cache (per core):  $\sim 4$  CPU cycles  
 L2 cache (per core pair):  $\sim 12$  CPU cycles  
 L3 cache/LLC (per uncore):  $\sim 28$  CPU cycles ( $\sim 25$  GiB/s)  
 DDR3-12800U RAM:  $\sim 28$  CPU cycles +  $\sim 50$ ns ( $\sim 12$  GiB/s)

## Caching – Cache Organisation

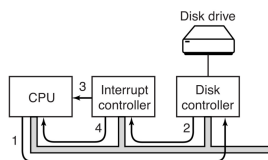
caches managed in hardware  
 divided into *cache lines* (usually 64 bytes each, unit at which data is exchanged between hierarchy levels)  
 often separation of data/instructions in faster caches (e.g. L1, see *harward architecture*)  
**cache hit**: accessed data already in cache (e.g. L2 cache hit)  
**cache miss**: accessed data has to be fetched from lower level  
 cache miss types:  
*compulsory miss*: first ref miss, data never been accessed  
*capacity miss*: cache not large enough for process working set  
*conflict miss*: cache has still space, but collisions due to placement strategy

## Interplay of CPU and Devices

I/O devices and CPU execute concurrently  
 Each device controller  
 - is in charge of particular device  
 - has local buffer

### Workflow:

1. CPU issues commands, moves data to devices
2. Device controller informs APIC that operation is finished
3. APIC signals CPU
4. CPU receives device/interrupt number from APIC, executes handler



## Device control

Devices controlled through their **device controller**, accepts commands from OS via **device driver**

devices controlled through device registers and device memory:  
*control* device by writing device registers  
*read* status of device by reading device registers  
*pass data* to device by reading/writing device memory

2 ways to access device registers/memory:

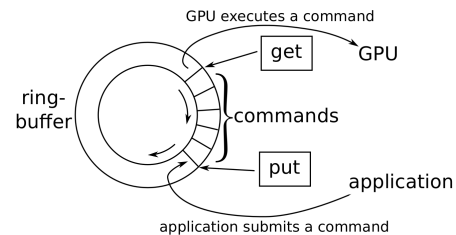
1. **port-mapped IO (PMIO)**:  
 use special CPU instructions to access port-mapped registers/memory  
 e.g. x86 has different *in/out*-commands that transfer 1, 2 or 4 bytes between CPU and device
2. **memory-mapped IO (MMIO)**:  
 use same address space for RAM and device memory  
 some addresses map to RAM, others to different devices  
 access device's memory region to access device registers/-memory

some devices use hybrid approaches using both

## Device control – Nvidia general purpose GPU

memory-mapped ring-buffer and *put/get*-device

mapping can be exposed to application  $\rightsquigarrow$  application can submit commands in user-mode



## Summary

The OS is an abstraction layer between applications and hardware (multiplexes hardware, hides hardware details, provides protection between processes/users)

The CPU provides a separation of User and Kernel mode (which are required for an OS to provide protection between applications)

CPU can execute commands faster than memory can deliver instructions/data – memory hierarchy mitigates this memory wall, needs to be carefully managed by OS to minimize slowdowns

device drivers control hardware devices through PMIO/MMIO

Devices can signal the CPU (and through the CPU notify the OS) through interrupts

### III. OS CONCEPTS

#### OS Invocation

OS Kernel does **not** always run in background!

Occasions invoking kernel, switching to kernel mode:

1. **System calls:** User-Mode processes require higher privileges
2. **Interrupts:** CPU-external device sends signal
3. **Exceptions:** CPU signals unexpected condition

#### System calls – motivation

Problem: protect processes from one another

Idea: Restrict processes by running them in user-mode

~> Problem: now processes cannot manage hardware,...  
who can switch between processes?  
who decides if process may open certain file?

~> Idea: OS provides **services** to apps  
app calls system if service is needed (**syscall**)  
OS checks if app is allowed to perform action  
if app may perform action and hasn't exceeded quota,  
OS performs action in behalf of app in kernel mode

#### System Calls – Examples

`fd = open(file, how, ...)` – open file for read/write/both

documented e.g. in `man 2 write`

overview in `man 2 syscalls`

#### System Calls vs. APIs

syscalls: interface between apps and OS services, limited number of well-defined entry points to kernel

APIs: often used by programmers to make syscalls  
e.g. `printf` library call uses `write` syscall

common APIs: Win32, POSIX, C API

#### System Calls – implementation

trap instruction: single syscall interface (entry point) to kernel  
switches CPU to kernel mode, enters kernel in same, predefined way for all syscalls

system call dispatches then acts as syscall multiplexer

syscalls identified by number passed to trap instruction

**syscall table** maps syscall numbers to kernel functions

dispatcher decides where to jump based on number and table

programs (e.g. `stdlib`) have syscall number compiled in!

~> never reuse old numbers in future kernel versions

#### Interrupts

devices use interrupts to signal predefined conditions to OS

reminder: device has „interrupt line“ to CPU

e.g. device controller informs CPU that operation is finished

**programmable interrupt controller** manages interrupts

interrupts can be **masked**

masked interrupts: queued, delivered when interrupt unmasked

queue has finite length ~> interrupts can get lost

notable interrupt examples:

1. *timer-interrupt:* periodically interrupts processes, switches to kernel ~> can then switch to different processes for fairness
2. *network interface card* interrupts CPU when packet was received ~> can deliver packet to process and free NIC buffer

when interrupted, CPU

1. looks up **interrupt vector** (= table pinned in memory, contains addresses of all service routines)
2. transfers control to respective **interrupt service routine** in OS that handles interrupt

interrupt service routine must first save interrupted process's state (instruction pointer, stack pointer, status word)

#### Exceptions

sometimes unusual condition makes it impossible for CPU to continue processing

~> **Exception** generated within CPU:

1. CPU interrupts program, gives kernel control
2. kernel determines reason for exception
3. if kernel can resolve problem ~> does so, continues **faulting instruction**
4. kills process if not

Difference to Interrupts: interrupts can happen in any context, exceptions always occur asynchronous and in process context

#### OS Concepts – Physical Memory

up to early 60s:

- programs loaded and run directly in *physical memory*
- program too large → partitioned manually into *overlays*
- OS then swaps overlays between disk and memory
- different jobs could observe/modify each other

#### OS Concepts – Address Spaces

bad programs/people need to be isolated

Idea: give every job the illusion of having all memory to itself  
every job has own *address space*, can't name addresses of others  
jobs always and only use virtual addresses

#### Virtual Memory – indirect addressing

Today: every CPU has built-in **memory management unit (MMU)**

MMU translates virtual addresses to physical addresses at every store/load operation

~> address translation protects one program from another

Definitions:

**Virtual address:** address in process' address space

**Physical address:** address of real memory

#### Virtual Memory – memory protection

MMU allows kernel-only virtual addresses

- kernel typically part of all address spaces
- ensures that apps can't touch kernel memory

MMU can enforce *read-only* virtual addresses

- allows safe sharing of memory between apps

MMU can enforce execute disable

- makes code injection attacks harder

#### Virtual Memory – page faults

not all addresses need to be mapped at all times

- MMU issues *page fault* exception when accessed virtual address isn't mapped
- OS handles page faults by loading faulting addresses and then continuing the program
- ~> memory can be **over-committed**: more memory than physically available can be allocated to application

page faults also issued by MMU on illegal memory accesses

#### OS Concepts – Processes

= program in execution („instance“ of program)

each process is associated with a **process control block (PCB)**  
contains information about allocated resources

each process is associated with a virtual **address space (AS)**

- all (virtual) memory locations a program can name
- starts at 0 and runs up to a maximum
- address 123 in AS1 generally ≠ address 123 in AS2
- indirect addressing ~> different ASes to different programs
- ~> protection between processes

## OS Concepts – address space layout

address spaces typically laid-out in different **sections**

- memory addresses between sections **illegal**
- illegal addresses  $\rightsquigarrow$  page fault
- more specifically called **segmentation fault**
- OS usually kills process causing segmentation fault

**Stack:** function history, local variables

**Data:** Constants, static/global variables, strings

**Text:** Program code

