

I. PROCESSES

The Process Abstraction

computers do „several things at the same time“ (just looks this way though quick process switching (**Multiprogramming**))

- ~> **process** abstraction models this concurrency:
 - container contains information about program execution
 - conceptually, every process has own „virtual CPU“
 - execution context is changed on process switch
 - dispatcher switches context when switching processes
 - **context switch**: dispatcher saves current registers/memory mappings, restores those of next process

Process-Cooking Analogon

Program/Process like Recipe/Cooking

Recipe: lists ingredients, gives algorithm what to do when

~> program describes memory layout/CPU instructions

Cooking: activity of using the recipe

~> process is activity of executing a program

multiple similar recipes for same dish

~> multiple programs may solve same problem

recipe can be cooked in different kitchens at the same time

~> program can be run on different CPUs at the same time (as different processes)

multiple people can cook one recipe

~> one process can have several worker threads

Concurrency vs. Parallelism

OS uses concurrency + parallelism to implement multiprogramming

1. **Concurrency**: multiple processes, one CPU
 - ~> not at the same time
2. **Parallelism**: multiple processes, multiple CPU
 - ~> at the same time

Virtual Memory Abstraction – Address Spaces

every process has own *virtual addresses* (vaddr)

MMU relocates each load/store to *physical memory* (pmem)

processes never see physical memory, can't access it directly

+ MMU can enforce protection (mappings in kernel mode)

+ programs can see more memory than available
80:20 rule: 80% of process memory idle, 20% active
can keep working set in RAM, rest on disk

- need special MMU hardware

Address Space (Process View)

code/data/state need to be organized within process

~> **address space layout**

Data types:

1. *fixed size* data items
2. data naturally *free'd in reverse allocation order*
3. data *allocated/free'd „randomly“*

compiler/architecture determine how large int is and what instructions are used in text section (**code**)

Loader determines based on exe file how executed program is placed in memory

Segments – Fixed-Size Data + Code

some data in programs never changes or will be written but never grows/shrinks

~> memory can be statically allocated on process creation

BSS segment (*block started by symbol*):

- statically allocated variables/non-initialized variables
- executable file typically contains starting address + size of BSS
- entire segment initially 0

Data segment:

- fixed-size, initlized data elements (e.g. global variables)

read-only data segment:

- constant numbers, strings

All three sometimes summarized as one segment

compiler and OS decide ultimately where to place which data/how many segments exist

Segments – Stack

some data naturally free'd in reverse allocation order

~> very easy memory management (stack grows upwards)

fixed segment starting point

store top of latest allocation in **stack pointer** (SP)
(initialized to starting point)

allocate **a** byte data structure: **SP += a; return(SP - a)**

free **a** byte data structure: **SP -= a**

Segments – Heap (Dynamic Memory Allocation)

some data „randomly“ allocated/free'd

two-tier memory allocation:

1. allocate large memory chunk (**heap segment**) from OS
 - base address + **break pointer** (BRK)
 - process can get more/give back memory from/to OS
2. dynamically partition chunk into smaller allocations
 - **malloc/free** can be used in random order
 - purely user-space, no need to contact kernel

